Human Face Generation using GAN

A major project report submitted in partial fulfillment of the requirement for the award of degree of

Bachelor of Technology

in

Computer Science & Engineering / Information Technology

Submitted by

Aniket Rawat (211315), Shiven Singh (211271), Uday Bhardwaj (211108)

Under the guidance & supervision of

Dr. Ramesh Narwal & Ms. Seema Rani

Department of Computer Science & Engineering and Information Technology

Jaypee University of Information Technology, Waknaghat, Solan - 173234 (India)

May 2025

SUPERVISOR'S CERTIFICATE

This is to certify that the major project report entitled **'Human Face Generation using GAN'**, submitted in partial fulfillment of the requirements for the award of the degree of **Bachelor of Technology** in **Computer Science & Engineering**, in the Department of Computer Science & Engineering and Information Technology, Jaypee University of Information Technology, Waknaghat, is a bona fide project work carried out under my supervision during the period from January 2025 to May 2025.

I have personally supervised the research work and confirm that it meets the standards required for submission. The project work has been conducted in accordance with ethical guidelines, and the matter embodied in the report has not been submitted elsewhere for the award of any other degree or diploma.

10/05/2025
(Supervisor Signature)
Supervisor Name: Dr. Ramesh Narwal
Designation: Assistant Professor
Department: Dept. of CSE & IT

16/05/2025
(Supervisor Signature)
Supervisor Name: Ms. Seema Rani
Designation: Assistant Professor
Department: Dept. of CSE & IT

Date: 8th May 2025
Place: Jaypee University of Information Technology, Solan

CANDIDATE'S DECLARATION

We hereby declare that the work presented in this report entitled **'Human Face Generation using GAN'** in partial fulfillment of the requirements for the award of the degree of **Bachelor of Technology** in **Computer Science & Engineering / Information Technology** submitted in the Department of Computer Science & Engineering and Information Technology, Jaypee University of Information Technology, Waknaghat is an authentic record of my own work carried out over a period from January 2025 to May 2025 under the supervision of **Dr. Ramesh Narwal & Ms. Seema Rani**.

The matter embodied in the report has not been submitted for the award of any other degree or diploma.

(Student Signature)
Name: Aniket Rawat
Roll No.: 211315
Date: 8th May 2025

(Student Signature)
Name: Shiven Singh
Roll No.: 211271
Date: 8th May 2025

(Student Signature)
Name: Uday Bhardwaj
Roll No.: 211108
Date: 8th May 2025

This is to certify that the above statement made by the candidates is true to the best of my knowledge.

10/05/2025
(Supervisor Signature with Date)
Supervisor Name: Dr. Ramesh Narwal
Designation: Assistant Professor
Department: Dept. of CSE & IT
Date: 8th May 2025

Verma
10/05/25
(Supervisor Signature with Date)
Supervisor Name: Ms. Seema Rani
Designation: Assistant Professor
Department: Dept. of CSE & IT
Date: 8th May 2025

ACKNOWLEDGEMENT

Firstly, I express my heartiest thanks and gratefulness to almighty God for His divine blessing makes it possible to complete the project work successfully.

I am really grateful and wish my profound indebtedness to Supervisors **Dr. Ramesh Narwal(Assistant Professor) & Ms. Seema Rani(Assistant Professor)**, Department of CSE, Jaypee University of Information Technology, Wakhnaghat. Deep Knowledge & keen interest of my supervisors in the field of **Artificial Intelligence** to carry out this project. Their endless patience, scholarly guidance, continual encouragement, constant and energetic supervision, constructive criticism, valuable advice, reading many inferior drafts and correcting them at all stages have made it possible to complete this project.

I would like to express my heartiest gratitude to **Dr. Ramesh Narwal & Ms. Seema Rani**, Department of CSE, for their kind help to finish my project.

I would also generously welcome each one of those individuals who have helped me straightforwardly or in a roundabout way in making this project a win. In this unique situation, I might want to thank the various staff individuals, both educating and non-instructing, which have developed their convenient help and facilitated my undertaking.

Finally, I must acknowledge with due respect the constant support and patients of my parents.

Aniket Rawat (211315)
Shiven Singh (211271)
Uday Bhardwaj (211108)

TABLE OF CONTENTS

LIST OF FIGURES

ABSTRACT

The attempt in the project “Human Face Generation Using GAN (Generative Adversarial Network)” is the attempt to replicate and imitate human faces using deep learning. This involves leveraging GANs, a state-of-the-art neural network framework consisting of two networks: there is a generator and a discriminator. The generator reproduces real faces during image generation while discriminator distinguishes the realistic ones, thus the two promoting healthy competition.

The main idea of the project is to condition the GAN on human faces’ data set to allow the network to learn and mimic real human faces attributes such as the face expression, the texture, and the symmetry of faces. This work proves that GAN can be proficient for image synthesis tasks, for example, creating realistic facial images for art professions, virtual reality and game industry.

Some of the issues targeted include; keeping face diversity, mode collapse and generating high resolution images. The present outcomes of the project reveal the model to create realistic-looking human faces which are difficult to distinguish from actual ones and underline the role of the GAN for the development of computer vision and artificial intelligence.

CHAPTER 1: INTRODUCTION

1.1 INTRODUCTION

Synthetic human face generation has rapidly evolved from early rule-based and statistical approaches to deep learning methods capable of producing photorealistic and highly diverse outputs. However, most existing systems take either purely generative inputs (random noise) or require an existing image to condition on, limiting their applicability in scenarios where only a textual description is available. Recent advances in cross-modal models—combining OpenAI’s CLIP for text-to-image alignment, StyleGAN2-ADA for high-fidelity generation, and ControlNet for precise structural control—offer a novel pipeline: convert a natural language prompt into a sketch and then produce a final, defect-free mugshot-style image.

At the heart of this pipeline is an adversarial training loop. A sketch generator, guided by CLIP embeddings extracted from the input text, produces line-drawing approximations of described facial features. A discriminator network then learns to distinguish these AI-generated sketches from hand-drawn forensic sketches, pushing the generator to improve realism and fidelity. The second phase exploits ControlNet's architecture to incorporate both the produced sketch and original text prompt, conditioning a StyleGAN2-ADA model to complete skin texture, lighting, and minute details. Refining between discriminator and generator iteratively produces images that are essentially unidentifiable from actual mugshots.

This project addresses two core domains. First, it broadens the use of GANs to forensic sketching, in which law enforcement depends on witness accounts instead of images. Second, it illustrates how text, structure, and generative priors can be combined for more controllable image synthesis. Beyond its technical advancements, the research instigates important debates about the ethical use of deep-fake enabled technologies in law enforcement and privacy contexts.

1.2 PROBLEM STATEMENT

The process of producing realistic human faces from GANs is a challenge and an issue:

1. **High Computational Overhead:**Two-stage pipeline training—the generation of a sketch and high-resolution face synthesis—is a significant draw on GPU resources and time, especially while fine-tuning StyleGAN2-ADA under ControlNet conditioning.

2. **Cross-Modal Alignment Challenges:** Ensuring that the sketch produced from a text prompt accurately captures subtle descriptors (e.g., "arched brows," "square jawline," "thin lips") requires strong CLIP-based encodings and sensitive loss design to prevent semantic drift.

3. **Adversarial Training Instability:** Both the sketch generator/discriminator and the final image generator/discriminator pairs can suffer from mode collapse, vanishing gradients, or oscillatory behavior, making convergence difficult without expert hyperparameter tuning.

4. **Structural Consistency vs. Photorealism:** ControlNet must balance adherence to the input sketch (to preserve witness-provided structural cues) with the GAN's tendency to hallucinate textures. Over-reliance on either can yield distorted features or overly generic, "blurry" results.

5. **Data Bias and Generalization:** Forensic cases span diverse demographics, yet training data is often skewed toward certain ethnicities, ages, or lighting conditions. The system must generalize to underrepresented groups to avoid reinforcing biases.

6. **Ethical and Privacy Considerations:** Automating sketch-to-mugshot pipelines raises concerns around misuse in creating deep-fakes, as well as the legal

admissibility and privacy implications of synthetic facial images. Robust safeguards and transparent reporting are essential to mitigate potential harms.

Solving these issues is essential for obtaining accurate, high-quality results and for determining the important contexts of this technology.

1.3 OBJECTIVES

The primary objectives of this project are:

1. **Develop a Two- Stage GAN Pipeline:** Design and apply a sketch- to- image system that first generates a facial sketch from textbook prompts using an OpenAI ‐ CLIP – conditioned creator and also produces a high- quality mugshot via a ControlNet – stoked StyleGAN2- ADA model.

2. **Achieve Photorealistic Detail:** Optimize the final generation stage to capture fine- granulated facial features similar as skin pores, subtle lighting variations, hair beaches, and emotion-driven expressions — at judgments suitable for forensic use.

3. **Ensure Cross-Demographic Diversity:** Enable the channel to synthesize accurate representations of all periods, genders, and races, as well as a full range of facial expressions, minimizing bias in sketch interpretation and final image literalism.

4. **Alleviate GAN- Specific Challenges:** It is easier to develop strategies for handling some of the challenges that are associated with GANs:
   - Training Insecurity: Employ adaptive literacy schedules, grade penalties, and progressive addition to stabilize inimical training across both sketch and final- image networks.

- Mode Collapse: Integrate diversity- promoting losses and expansive data addition to help repetitious labors.
- Computational Effectiveness: Influence mixed-perfection training and model pruning to reduce GPU time without demeaning affair quality.

5. **Explore Forensic and marketable operations:** Probe practical deployments in law enforcement(automated suspect mugshots), entertainment(custom icon creation), and data-sequestration surrounds(synthetic datasets for training vision models without real individualities).

Through these objectives, the project aims to balance technical innovation with responsible development, providing practical outputs and insights into the transformative potential of GANs.

1.4 SIGNIFICANCE AND MOTIVATION OF THE PROJECT WORK

This design holds both specialized significance and real-world impact:

1. **Forensic Innovation:** Automating the conversion of substantiation descriptions into high- dedication mugshots can accelerate felonious examinations while reducing homemade sketch artist workload, potentially leading to briskly questionable identification.

2. **Advancement inCross-Modal Generation:** By integrating CLIP, StyleGAN2-ADA, and ControlNet, the work pushes the boundaries of textbook- to- image conflation exploration, demonstrating how structural guidance can upgrade photorealistic labors.

3. **Synthetic Data for sequestration-Sensitive disciplines:** Producing different, realistic facial images without counting on real individualities supports

sequestration- conserving datasets for training biometric and surveillance models in healthcare, security, and academic exploration.

4. **Creative and marketable operations:** Beyond forensics, the channel can drive cost-effective icon creation for gaming, virtual reality, and digital media, offering substantiated characters from simple textual inputs.

5. **Ethical mindfulness and Safeguards:** Pressing both the pledge and threats of deep-fake-able systems, the design motivates the development of countermeasures similar as bedded watermarks and forensic discovery tools to insure responsible deployment.

6. **Educational Value:** Experimenters and interpreters gain hands- on experience with inimical training,multi-modal exertion, and large- scale dataset curation, heightening understanding of ultramodern GAN infrastructures and their optimization challenges.

The provocation for this work stems from the binary need to harness the creative power of GANs for salutary operations(e.g., forensic sketching, icon design) while proactively addressing the societal pitfalls posed by unbridled synthetic face generation.

1.5 ORGANIZATION OF PROJECT REPORT

To ensure a structured approach, this report is divided into the following chapters:

1. **Introduction:** Introduces the project, describes what the work is about, formulates the problem statement, defines objectives, and explains why the study is of value.

2. **Literature Survey:** Summarises prior work conducted in the area of GAN-based face generation, outlines some findings of the project, and discusses some of the research questions and holes in the present literature that this project aims to fill.

3. **System Development:** Explains how the system could be designed to meet the identified requirements, how data would be pre-processed, how various implementation strategies may be applied, and relates any emergent issues and problems encountered.

4. **Testing:** Presents the testing techniques used to check the authenticity of the system, provides test examples, and outlines test results.

5. **Results and Evaluation:** Discusses findings of the project, assesses their quality against set standards and benchmark existing solutions in analyses where relevant.

6. **Conclusions and Future Scope:** Repeats the main conclusions derived from the project, admits its shortcomings, and outlines lines of future research and development.

This report should help the reader gain knowledge about the project from the problem definition stage up to the evaluation of the results and their significance.

CHAPTER 2: LITERATURE SURVEY

2.1 OVERVIEW OF RELEVANT LITERATURE

Title	Author	Year of Publication	Publishing Details	Summary
LumiGAN: Unconditional Generation of Relightable	Boyang Deng, Yifan Wang, Gordon Wetzstein	2023	arXiv:2304.13153	Introduces a GAN that generates relightable, geometry-accurate human faces, achieving new state-of-the-art FID on FFHQ-256 and FFHQ-1024 by combining photorealism with precise visibility predictions. Unlike earlier 3D GANs, it balances anatomical correctness with geometric fidelity, setting a benchmark for relightable face synthesis.
StyleNAT: Giving Each Head a New Perspective	Steven Walton, Ali Hassani, Xingqian Xu, Zhangyang Wang, Humphrey Shi	2022	arXiv:2211.05770	Proposes a lightweight generative architecture that matches top visual quality on FFHQ-256 and FFHQ-1024 while drastically reducing parameters and boosting sampling speed. It prevents extreme latent-code artifacts, ensuring smooth style transitions and efficient image production.
Reconstruct Face from	Xingbo Dong, Zhihui Miao,	2022	arXiv:2206.04295	Frames facial reconstruction as a

Features Using GAN Generator as a Distribution Constraint	Lan Ma, Jiajun Shen, Zhe Jin, Zhenhua Guo, Andrew Beng Jin Teoh			GAN-constrained optimization to recover full faces from partial inputs, demonstrating 98.0% LFW-dataset recovery under type-I attacks at 0.1% FAR. The technique remains resilient in adversarial situations and promises applications for secure surveillance and biometric systems.
Fake Face Generator: Generating Fake Human Faces using GAN	Md. Mahiuddin , Md. Khaliluzzaman , Md. Shahnur Azad Chowdhury, Muhammed Nazmul Arefin	2022	International Journal of Advanced Computer Science and Applications, Vol. 13, No. 7	Describes a generator–discriminator model that completes lost face areas using thin features and maximizes adversarial losses for producing photorealistic faces without relying on auxiliary classifiers. The solution prioritizes computational efficiency without sacrificing image quality for real-world deployment
AniGAN: Style-Guided Generative Adversarial Networks	Bing Li, Yuanlue Zhu, Yitong Wang, Chia-Wen Lin, Bernard Ghanem, Linlin Shen	2022	IEEE Transactions on Multimedia, vol. 24	Creates a style-guided GAN that generates artifact-free anime portraits that are in line with reference styles and surpasses state-of-the-art on selfie2anime and face2anime datasets. Its style-accurate, smooth results render it suitable for gaming, animation, and digital art where stylistic control is paramount.

BlendGAN: Implicitly GAN Blending for Arbitrary StylizedFace Generation	Mingcong Liu, Qiang Li, Zekui Qin, Guoxin Zhang, Pengfei Wan, Wen Zheng	2021	arXiv:2110.11728	Introduces a GAN that combines various artistic styles to produce expressive faces, supported by a weight-decoupled module (WDM) on the AAHQ dataset. Although it performs well in style consistency, processing some reference styles is still an issue to be addressed in future refinement.
Artificial (or) Fake Human Face Generator using Generative Adversarial Network (GAN) Machine Learning Model	Shariff, Daanish Mohammed, H. Abhishek, and D. Akash	2021	Fourth International Conference on Electrical, Computer and Communication Technologies (ICECCT)	Employs a traditional DCGAN to generate CelebA images, with moderate structural correspondence (0.34) to originals, highlighting the baseline ability of the model. While training is computationally demanding, it provides baseline insights into photorealistic head generation based on GANs.
FaceShifter: Towards High Fidelity AndOcclusion AwareFace Swapping	Lingzhi Li, Jianmin Bao, Hao Yang, Dong Chen, Fang Wen	2020	arXiv:1912.13457	Introduces a face-swapping GAN that delivers high-quality, occlusion-robust transformations on CelebA-HQ and FFHQ, surpassing previous methods in realism under sunglasses and varied poses. Infrequent blur and loss of detail identify regions for improving edge acuity and attribute retention.

Interpreting the Latent Space of GANs for Semantic Face Editing	Shen, Y., Gu, J., Tang, X., & Zhou, B	2020	CVPR 2020	Based on PGGAN/StyleGAN to support fine-grained control over attributes such as age and expression, tested on CelebA-HQ and FFHQ with realistic realism. Excessive latent manipulations may still create unnatural artifacts, which indicates the necessity for boundary-conscious editing constraints
ArcFace: Additive Angular Margin Loss forDeep Face Recognition	Deng, Jiankang, Jia Guo, Niannan Xue, and Stefanos Zafeiriou	2019	Journal of Latex Class files, vol. 14, no. 8	Introduces an additive angular margin loss in a DCNN that is shown to drastically enhance face identification performance on ten benchmarks of MS1MV2, promoting discriminative feature learning. The research is centered on recognition performance and does not cover more general integration or image synthesis tasks.
Wav2Pix: Speechconditi on ed F ace Generation using Generative Adversarial Networks	Duarte, A.C., Roldan, F., Tubau, M., Escur, J., Pascual, S., Salvador	2019	ICASSP 2019	Suggests a GAN which converts speech intervals into matching facial imagery, conditioned on proprietary YouTube data, demonstrating proof-of-concept for audio-conditioned image synthesis. Limitations are low-resolution outputs and sensitivity to short audio clips, suggesting scope for

				scalability and noise robustness.
A Style-Based Generator Architecture for Generative Adversarial Networks	Tero Karras, Samuli Laine, Timo Aila	2019	CVPR 2019	Describes a style-based GAN using Adaptive Instance Normalization to decouple content and style, outperforming classical GANs on CelebA-HQ and FFHQ in variability and output quality. Its latent space constraints still entangle style and content, suggesting further work to enhance disentanglement.

2.2 Key Gaps in the Literature

1. **LumiGAN: Unconditional Generation of Relightable**

   Entirely limited to face generation; It does not support dynamic scene generation such as relighting or animating ability required for interactive applications.

2. **StyleNAT: GivingEach Head a New Perspective**

   Problems with latent space translation because of the shift in the latent code - moving it causes the generation of unrealistic images and limits the results to extremes.

3. **Reconstruct Face from Features Using GAN Generator as a Distribution Constraint**

   Concerned with attack types, but does not consider measures to avoid using face reconstruction for threats to security.

4. **Fake Face Generator: Generating Fake Human Faces using GAN**

   The model lacks diversification in the generated faces and can only generate pictures of limited resolution. Mostly limited in terms of falsification to simple fake face creation.

5. **AniGAN: Style-Guided Generative Adversarial Networks**

   Single dataset testing (selfie2anime and face2anime) and difficulties in extending the method and using it for different, non-anime examples.

6. **BlendGAN: Implicitly GAN Blending for Arbitrary StylizedFace Generation**

   Not compatible with some of the reference style types reflected in the AAHQ dataset, thus, not applicable to several more styles.

7. **Artificial (or) Fake Human Face Generator using Generative Adversarial Network (GAN) Machine Learning Model**

   Hence, moderate structural similarity, (MSS) of 0.34, is an indication that limitations exist in generating high fidelity output in using the approach and does not show a marked enhancement in efficiency for the construction of virtual environments.

8. **FaceShifter: Towards High Fidelity AndOcclusion AwareFace Swapping**

   Becomes involved with problems such as vagueness of artifacts and loss of some attributes, which makes it difficult to get realistic in complex circumstances.

9. **Interpreting the Latent Space of GANs for Semantic Face Editing**

   High levels of latent space manipulation generate unrealistic changes, proving the minimum realism for larger-scale EAs.

10. **ArcFace: Additive Angular Margin Loss forDeep Face Recognition**

    Largely concerned with relative performance and not with the actual integration of the system into a more extensive system environment.

11. **Wav2Pix: Speech Conditioned Face Generation using Generative Adversarial Networks**

    It declines greatly with smaller speech chunks or low definition images, raising questions on scalability and robustness of the system.

12. **A Style-Based Generator Architecture for Generative Adversarial Networks**

    However, based on the disentanglement results , there are always entangled constraints on the distribution of the latent space (Z), which inevitably cause entanglement when separating style and content from the training data.

Most current GAN-based face-generation frameworks struggle when the latent vectors are pushed to extreme values, often producing distorted or unrealistic images. Their architectures are typically fine-tuned for specific datasets, limiting transferability to more diverse, real-world scenarios where varied lighting, poses, and demographics are present. Efforts to scale resolution often fail above modest sizes, limiting models to simple, static environments. Traditional artifacts—noise patterns, blurring, unnatural textures—still prove hard to remove, challenging the robustness and resilience of these systems. In addition, ethical concerns related to deep-fakes and abuse are rarely part of the model design, making mitigation strategies immature. Lastly, there is a significant lack of proven real-world deployments or integration roadmaps, i.e., that high-level system-level considerations—ranging from data pipelines to legal compliance—are left largely unexamined.

CHAPTER 3: SYSTEM DEVELOPMENT

3.1 REQUIREMENTS AND ANALYSIS

This part gives a technical overview of the system requirements such as functional and non-functional requirements, the technology ecosystem, hardware dependencies, and dataset considerations vital to the successful implementation of an efficient and scalable solution for producing human faces based on GANs.

1. **Functional Requirements**
   - **Text-to-Sketch Conversion:** Accept a text description and generate a legible, structurally correct outline of a human face through an OpenAI CLIP–conditioned generator.The model should correctly convert a natural language description into a sketch image of a human face. This is done by a CLIP-conditioned generator in which text semantic features are translated to face structure and rendered in the form of grayscale line sketches. These sketches should have high levels of detail such as facial shape, hair, and distinguishing facial features useful for forensic realism.

   - **Sketch-to-Mugshot Synthesis:** Take the generated sketch and the original text prompt to yield a high-resolution, photorealistic “mugshot” image through a ControlNet–guided StyleGAN2-ADA pipeline.Using the generated sketch and the original text as dual conditioning, the system must produce a **photorealistic, high-resolution face**. This is handled by the **ControlNet-enhanced StyleGAN2-ADA model**, which completes texture synthesis, lighting realism, and fine-grain expression control. The result must be comparable to real human portraits.

- **Alignment Validation:** Compute edge-based alignment scores between sketch and final image to iteratively refine generation until a predefined threshold (e.g. > 0.75) is reached.The model must iteratively refine outputs using **structural consistency metrics**. A custom alignment score compares sketch edges and final image features. Generation loops until an edge overlap threshold (≥0.75) is reached, ensuring the result is structurally faithful to the sketch and by extension, the input prompt.

2. Non-Functional Requirements

In addition to feature-level expectations, the system must satisfy broader operational requirements to ensure reliability and maintainability:

- **Stability & Scalability:** Ensure adversarial training (for any fine-tuning) remains stable, and that the inference pipeline handles batch processing of prompts without crashes.The system must avoid typical adversarial training issues like **mode collapse or vanishing gradients**, especially during sketch generation and GAN fine-tuning. It should also be capable of handling **batch prompt processing** during inference without performance degradation or crashes.

- **Performance:** Achieve end-to-end generation (sketch + mugshot) within a practical time frame (e.g. < 1 minute per prompt) on high-end GPUs.For usability in real-world or near-real-time applications, the **total runtime per prompt (sketch + mugshot)** should not exceed 60 seconds on a high-end GPU (such as an NVIDIA RTX 3090 or A100). Optimizations like float16 precision and attention slicing should be employed to meet this requirement.

- **Maintainability:** Organize code into modular sections—installation, imports, model loading, sketch generation, final synthesis, evaluation—to facilitate future extensions.The codebase should be modular and well-structured, enabling easy debugging, extension, and re-training. Each major function—from **model loading** to **final image evaluation**—should reside in independent modules for clarity and future development flexibility.

3. Technical Requirements

To ensure smooth deployment, specific tools, libraries, and environments are required:

- **Programming Language & Frameworks:** Python 3.8+, PyTorch, NVLabs' StyleGAN2-ADA repository, Diffusers, Transformers.

- **Hardware:** NVIDIA GPU with CUDA 11+ (e.g. RTX 3090 or A100) and ample VRAM for half-precision (float16) inference.

- **Libraries:** torch, dnnlib, legacy (StyleGAN2-ADA), clip, diffusers, opencv-python, PIL, numpy.

- **Supporting Tools:** CUDA toolkit, ninja, Git for cloning and version control.

4. Dataset Analysis

- **Primary Image Corpus:** FFHQ (Flickr-Faces-HQ) at 1024×1024 resolution, providing diverse, high-quality face images for StyleGAN2-ADA.The system uses the **FFHQ (Flickr-Faces-HQ)** dataset which contains over 70,000 high-quality human face images at 1024×1024 resolution. This dataset is favored for its wide demographic coverage, including age, gender, skin tone, and facial geometry, making it suitable for training generative models with high generalization ability.

- **Sketch Supervision:** Optional collection of hand-drawn forensic sketches to fine-tune edge-detector parameters; otherwise, multi-phase Canny + Sobel processing on FFHQ outputs.The initial face image produced by the GAN is converted to a sketch using **multi-phase edge detection** involving Canny and Sobel filters. In cases where access to **real forensic sketches** is available, the sketch generator can be optionally fine-tuned to better mimic forensic styles.

- **Preprocessing Challenges:** Harmonizing image formats and resolutions (cropping to 512×640 for sketch stage, 768×1024 for final synthesis), normalizing pixel intensities to [–1,1], and balancing demographic representation to mitigate bias.A few specific challenges arise while standardizing the dataset for input into the GAN and diffusion modules:

    - **Image Resizing**: FFHQ images are cropped or scaled to **512×640** for sketch input and **768×1024** for final synthesis.

    - **Intensity Normalization**: Pixel values are rescaled to [–1, 1] for StyleGAN2 and [0, 1] for Stable Diffusion.

    - **Bias Mitigation**: The FFHQ dataset, while diverse, still exhibits demographic imbalance. Careful selection and balancing of prompt–image pairs are necessary to prevent overfitting or biased outputs.

Figure 1. *FFHQ Dataset*

3.2 PROJECT DESIGN AND ARCHITECTURE

This section outlines the complete design blueprint of the system, encompassing both the **conceptual workflow** and the **technical architecture**. It explains how various deep learning models and modules are integrated to generate high-quality human faces from plain text descriptions.

1. **High-Level Overview**
    - **Text Encoding & Sketch Generation:** The process begins with the text prompt provided by the user. The system uses CLIP's vision transformer model ("ViT-B/32") to tokenize and encode this text prompt into a set of numerical features. However, to ensure that the generated face aligns with the desired emotional attributes, emotion-related features in the text prompt are masked out during encoding. The resulting features are then projected into a latent space, where they are mapped through StyleGAN2's mapping network to generate an initial face image.The initial face image is

subject to multi-phase edge detection processes, which involve both Canny and Sobel detectors. The output of this step is a detailed forensic-style sketch, which serves as a foundation for the final image synthesis. The sketch helps to retain structural consistency and guides the subsequent face synthesis steps.

- **ControlNet-Guided Face Synthesis:** With the initial sketch in hand, the next phase involves feeding both the sketch and the original text prompt into a **StableDiffusionControlNetPipeline**. This pipeline utilizes the "Realistic Vision V6.0" base model, optimized for photorealistic image synthesis. Some of the major features of this optimization are the **UniPCMultistepScheduler**, which stabilizes the diffusion process, and the memory-effective attention mechanism, which makes it possible for the model to manage the computational burden efficiently.

  This module improves the initial outline and utilizes the context derived from the initial text prompt to inform the synthesis of a final, realistic facial image. This involves several steps of inference, wherein the image is incrementally improved until the intended level of detail and realism is reached.

2. Architecture Details

- **Sketch Generator:**
  - ★ *Mapping Network:* The mapping network is responsible for converting the text prompt into a format that the StyleGAN2 model can utilize. In particular, it employs fully connected layers to project the features learned by the CLIP model (512-D features) into the latent space of StyleGAN2 (also 512-D). This conversion is necessary to make sure that the resulting face image is both

consistent with the text input and meets the structural constraints outlined in the sketch.

- ★ *Synthesis Network:* After the features are projected onto the latent space, they go through a synthesis network based on progressive convolutions. The network progressively refines the image by upsampling in several stages up to a resolution of 1024×1024 pixels. This high-resolution output helps in maintaining fine detail in the final image, something that is absolutely necessary for building photorealistic human faces. The noise_mode="const" ensures that structure in the image is maintained along the way while refining.

- ★ *Edge Extraction:* Used combined Canny and Sobel detectors, morphological dilation and closing to make lines thicker. Edge extraction is an important stage in the pipeline because it assists in creating the first sketch that directs the synthesis of the face. It employs a combination of Canny and Sobel edge detectors, which are used to emphasize the prominent contours of the image. Upon primary edge detection, the image goes through morphological processing, involving closing and dilation, to fill in the lines and make it a clearer and more defined sketch.

- ● **Final Image Pipeline:** The last image synthesis is the combination of a number of state-of-the-art deep learning models, such as the ControlNet module and the Diffusion backbone.

  - ★ *ControlNet Module:* ControlNet module uses the original sketch and conditions it on the input text prompt. "Scribble" conditioning network helps in making sure that the ultimate output has close similarity with the sketch and original text input. This conditioning plays an essential role in upholding structural coherence and helping in maintaining correspondence of generated face features with the description by the user.

★ *Diffusion Backbone:* Stable Diffusion v1.5 or "Realistic Vision" with 30–95 inference steps and guidance_scale 7.5–11.0. The diffusion model, either Stable Diffusion v1.5 or Realistic Vision, is used as the backbone to generate the final face image. It follows a step-wise procedure, and the inference steps can be between 30 and 95. The scale of guidance, which is normally between 7.5 and 11.0, is used to steer the output of the model in the direction of desired facial structure and look. This operation sequentially refines the image so that it gets more realistic with every step.

★ *Post-Processing:*Bilateral filtering and sharpening kernel for the enhancement of pores and fine details. Once the final image is produced, it passes through a number of post-processing steps aimed at making it as high-quality as possible. The final image goes through bilateral filtering, which is meant to smoothen the noise but keep the edges, and a sharpening kernel is employed to further sharpen fine details like pores and hair textures. The last processes play an essential role in refining the visual attractiveness and realism of the final produced image.

3. Model Flow

- **Prompt → CLIP Encoding → Latent z:**
  It starts from the user's text prompt that is encoded with CLIP. Features are then projected into StyleGAN2's latent space, producing a set of initial features (latent z).
- **Latent z → StyleGAN2 Mapping → Synthesis Image**

The latent z is fed into StyleGAN2's mapping network, which produces an initial image from the encoded prompt. This image is used as a starting point for refinement.

- **Image → Edge Detection → Sketch:**
  The initial image is then subjected to edge detection, yielding a detailed forensic sketch that retains the structural aspects of the face.

- **Sketch + Prompt → ControlNet → Diffusion Steps → Final Image:**
  The sketch, together with the initial text prompt, is passed to the ControlNet module that conditions the generation of the image. The steps of diffusion work on refining the image, continuously adding detail and photorealism.

- **Compute Edge Alignment; Loop Until Threshold:**
  The model keeps refining the image, verifying the alignment between the edges of the sketch and the final image. This loop runs until the alignment score is greater than or equal to the predefined threshold, with a high level of accuracy in the generated face.

4. Evaluation Metrics

- **Alignment Score:** The alignment score is a measure that describes how well the edges in the output image correspond to those in the input sketch. The better the structural consistency, the higher the score. The desired threshold for this score is above 0.75 so that the generated image stays true to the original sketch.

- **FID & IS:** The **Frechet Inception Distance (FID)** and **Inception Score (IS)** are often employed metrics to assess the photorealism and diversity of outputs. The FID score indicates how close the generated image distribution is to a reference set of real images, with lower scores representing greater realism. The IS score measures the diversity of the generated images, with higher scores reflecting greater diversity of outputs.

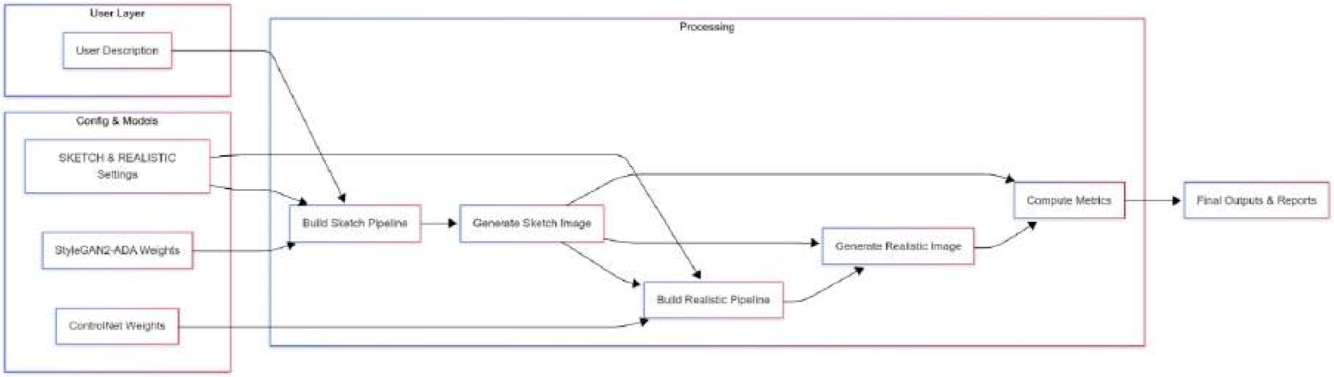

Figure 2. *Data Flow Diagram*

3.3 DATA PREPARATION

Data preparation is an important step in making sure that the system works properly. This phase includes preprocessing the images and text prompts to produce a uniform and high-quality input for the training and synthesis process. The objective of data preparation is to prepare the input data for the different deep learning models (StyleGAN2, Stable Diffusion, ControlNet) employed in the face generation system. Here, we detail the process as image preprocessing, text prompt processing, and training-validation split.

1. **Image Preprocessing :**Proper image preprocessing guarantees that the input data presented to the system is consistent, of high quality, and appropriate for training and inference. A number of steps constitute the image preprocessing pipeline:

- **Resolution Standardization:** All FFHQ images are first resized to a default resolution of 1024×1024 pixels. This is an ideal resolution for input to the StyleGAN2 model, which is meant to handle high-resolution images. Using a uniform image size facilitates better learning of facial feature representations by the model during training.

StyleGAN2 Input: The images are resized to 1024×1024 pixels prior to input into the network.

Subsequent Cropping/Skewing: Based on the individual requirements of each module, the images are cropped or skewed further to other resolutions. For example, certain modules might need a resolution of 512×640 or 768×1024, and the images are cropped accordingly to fulfill these demands.

This multi-resolution approach allows the system to maintain flexibility while ensuring that each module works with an appropriate input size for optimal performance.

- **Intensity Normalization:** For maintaining consistency of image values between models, the pixel intensities are normalized as per each deep learning model to be utilized within the system:
  For StyleGAN2: The pixel values are normalized to the range of [–1, 1], since StyleGAN2's structure accepts input images with pixel values within this interval.
  For Diffusion Models: The pixel values are normalized to the range [0, 1] to be compatible with the diffusion model, which usually operates on pixel values in this range for improved stability during training and inference.
  This normalization helps ensure that the models can handle the images properly and consistently, resulting in better-quality outputs.

- **Edge Enhancement:** To create sharp and accurate sketches from the original images, various edge enhancement methods are used:
  Canny Edge Detection: Canny edge thresholds are adjusted between 25 and 75 to detect prominent edges in the image. This process aids in

detecting prominent structural features of the face, which are significant in creating detailed sketches.

Sobel Filtering: Sobel operators are used to improve edge detection and further sharpen the contours of the image. Sobel filtering enhances edge detection by computing the gradient of the image intensity.

Morphological Operations: Following edge detection, morphological operations like dilation and closing are employed to broaden the lines so that they are thicker and more prominent, thereby appropriate for generating sketches.

These preprocessing steps make the final sketches clear and detailed, which enables them to be better aligned with the generated final image.

2. **Text Prompt Processing :** The quality of generated face images significantly relies on input text prompts, and therefore, it is highly important to preprocess these prompts in order to guarantee consistency and relevance in the output of the model.

- **Tokenization & Embedding:** The text prompts are initially tokenized with the CLIP tokenization scheme. This tokenization transforms the text into a format that is compatible with the deep learning models. Enforced prompt templates are employed to ensure consistency and avoid generating irrelevant variations.

  Instruction Template: A uniform template is imposed on every prompt, in the format of:
  "*security camera picture, {prompt}, front side, neutral expression, realistic illumination*."
  This template guarantees that every text prompt contains standard descriptors like the front view, neutral pose, and realistic lighting, which are essential for producing a consistent and realistic face. The addition of the term "security camera photo" promotes the production of mugshot-like images, giving further control over the style and composition.

- **Feature Masking:** As emotional expressions and features can widely differ based on the input prompt, feature masking is employed to suppress emotional aspects in generated faces. Feature masking is accomplished by a learned mask on the 512-D output of CLIP. The mask selectively inhibits some emotional features from the text prompt, so that face generation tends towards a neutral mugshot style instead of an expressive or varied face.

  This step ensures that the created faces are more aligned with the desired "mugshot" appearance, something especially important when the system needs to create faces for applications demanding neutral and objective portrayals.

3. **Training–Validation Split :** The training and validation process aims to make the model generalize and perform well on unseen data. Data split plays an important role in assessing how well the model performs on tasks like sketch fidelity, alignment, and demographic diversity.

   - **FFHQ Partitioning:** The FFHQ dataset, with a vast number of high-quality human face images, is split into two halves for training and validation: Training Set: 90% of the images are utilized for training, specifically for fine-tuning the sketch generator. This enables the system to learn to produce high-quality sketches that reflect the structural details of faces.
     Validation Set: The last 10% of the images are reserved for validation. These images are utilized to assess the accuracy of the sketches and verify that the generator is keeping high accuracy and diversity when generating sketches of unknown faces.By splitting the dataset in this way, the model can learn and be tested on how well it can generalize to new, unseen data.

- **Prompt Sets:** Besides partitioning the FFHQ images, a varied collection of 1,000 text prompts is selected so that they address a broad array of facial traits, poses, and lighting environments. The text prompts are grouped into two subsets:

  Training Set (80%): The bulk of the text prompts (80%) is applied to train the system and calibrate the face generation models. This helps the system learn to accommodate an extensive range of prompts and produce appropriate faces given various textual inputs.

  Testing Set (20%): The remaining 20% of the prompts are used for testing the final synthesis quality and evaluating the alignment performance. These prompts help assess how well the system can generalize to new and diverse prompts, as well as how accurately it can generate faces that match the textual descriptions.

  By using this split, the system can be rigorously tested for both the quality of the final images and its ability to align generated faces with the given text prompts.

3.4 IMPLEMENTATION

1. Text-to-Latent Projection

Define ‘text_to_latent(prompt: str) → torch.Tensor’ to encode a text prompt into a 512-D latent vector for StyleGAN2. It tokenizes the prompt (with enforced forensic-photo context), obtains CLIP embeddings, applies an emotion-suppression mask, and projects via a learnable linear layer.

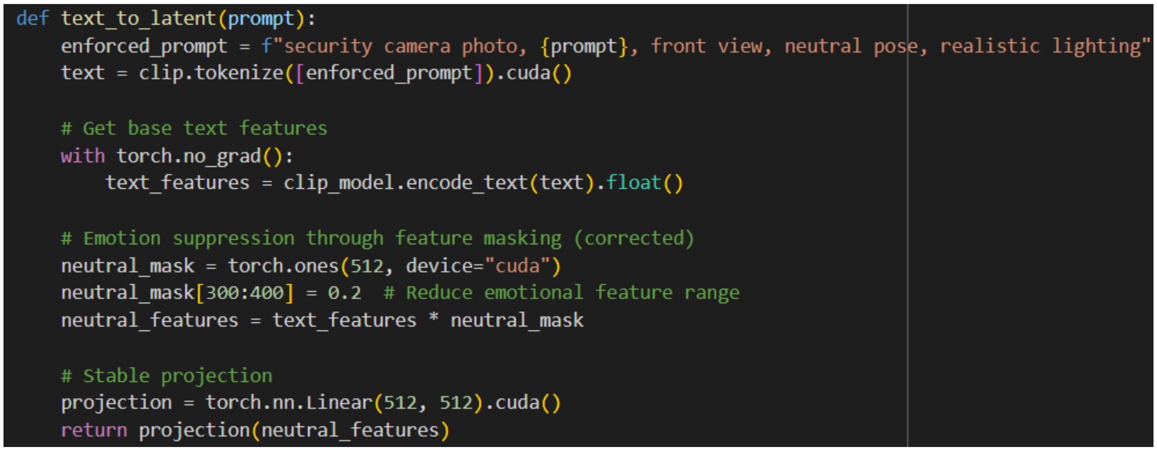

```python
def text_to_latent(prompt):
    enforced_prompt = f"security camera photo, {prompt}, front view, neutral pose, realistic lighting"
    text = clip.tokenize([enforced_prompt]).cuda()

    # Get base text features
    with torch.no_grad():
        text_features = clip_model.encode_text(text).float()

    # Emotion suppression through feature masking (corrected)
    neutral_mask = torch.ones(512, device="cuda")
    neutral_mask[300:400] = 0.2  # Reduce emotional feature range
    neutral_features = text_features * neutral_mask

    # Stable projection
    projection = torch.nn.Linear(512, 512).cuda()
    return projection(neutral_features)
```

Figure 3. *Workflow of CLIP-based text encoding with feature masking and linear projection.*

2. Sketch Generation via StyleGAN2

Implement ‘generate_sketch(prompt: str) → PIL.Image’ that maps the latent code through StyleGAN2’s mapping and synthesis networks, then extracts edges to form a forensic sketch. Multi-phase edge detection (Canny + Sobel) and morphological operations yield crisp line art.

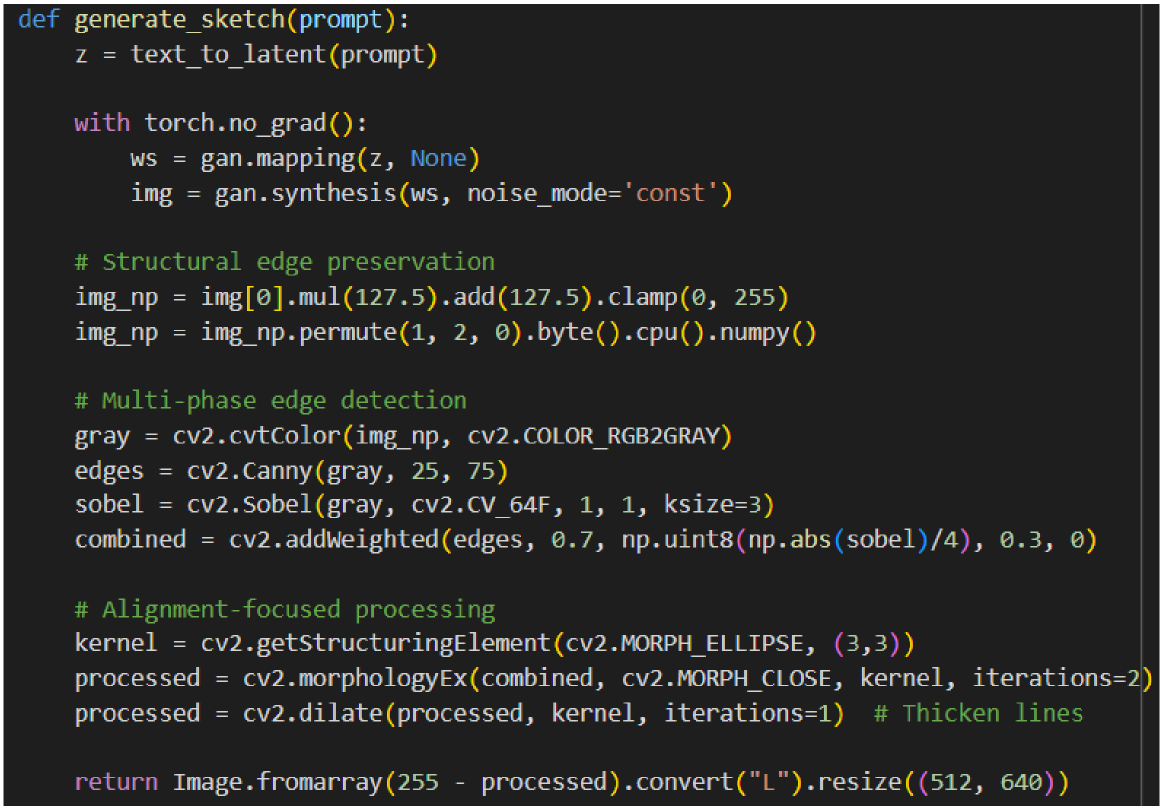

```
def generate_sketch(prompt):
    z = text_to_latent(prompt)

    with torch.no_grad():
        ws = gan.mapping(z, None)
        img = gan.synthesis(ws, noise_mode='const')

    # Structural edge preservation
    img_np = img[0].mul(127.5).add(127.5).clamp(0, 255)
    img_np = img_np.permute(1, 2, 0).byte().cpu().numpy()

    # Multi-phase edge detection
    gray = cv2.cvtColor(img_np, cv2.COLOR_RGB2GRAY)
    edges = cv2.Canny(gray, 25, 75)
    sobel = cv2.Sobel(gray, cv2.CV_64F, 1, 1, ksize=3)
    combined = cv2.addWeighted(edges, 0.7, np.uint8(np.abs(sobel)/4), 0.3, 0)

    # Alignment-focused processing
    kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (3,3))
    processed = cv2.morphologyEx(combined, cv2.MORPH_CLOSE, kernel, iterations=2)
    processed = cv2.dilate(processed, kernel, iterations=1)  # Thicken lines

    return Image.fromarray(255 - processed).convert("L").resize((512, 640))
```

Figure 4. *Pipeline for converting a synthesized face into a refined sketch via edge detection and morphology.*

3. Final Mugshot Synthesis with ControlNet

The function ‘generate_final_image(sketch: PIL.Image, prompt: str) → PIL.Image’ loads a ControlNet-augmented Stable Diffusion pipeline, runs conditioned diffusion over 30–60 steps, then applies bilateral filtering for realism.

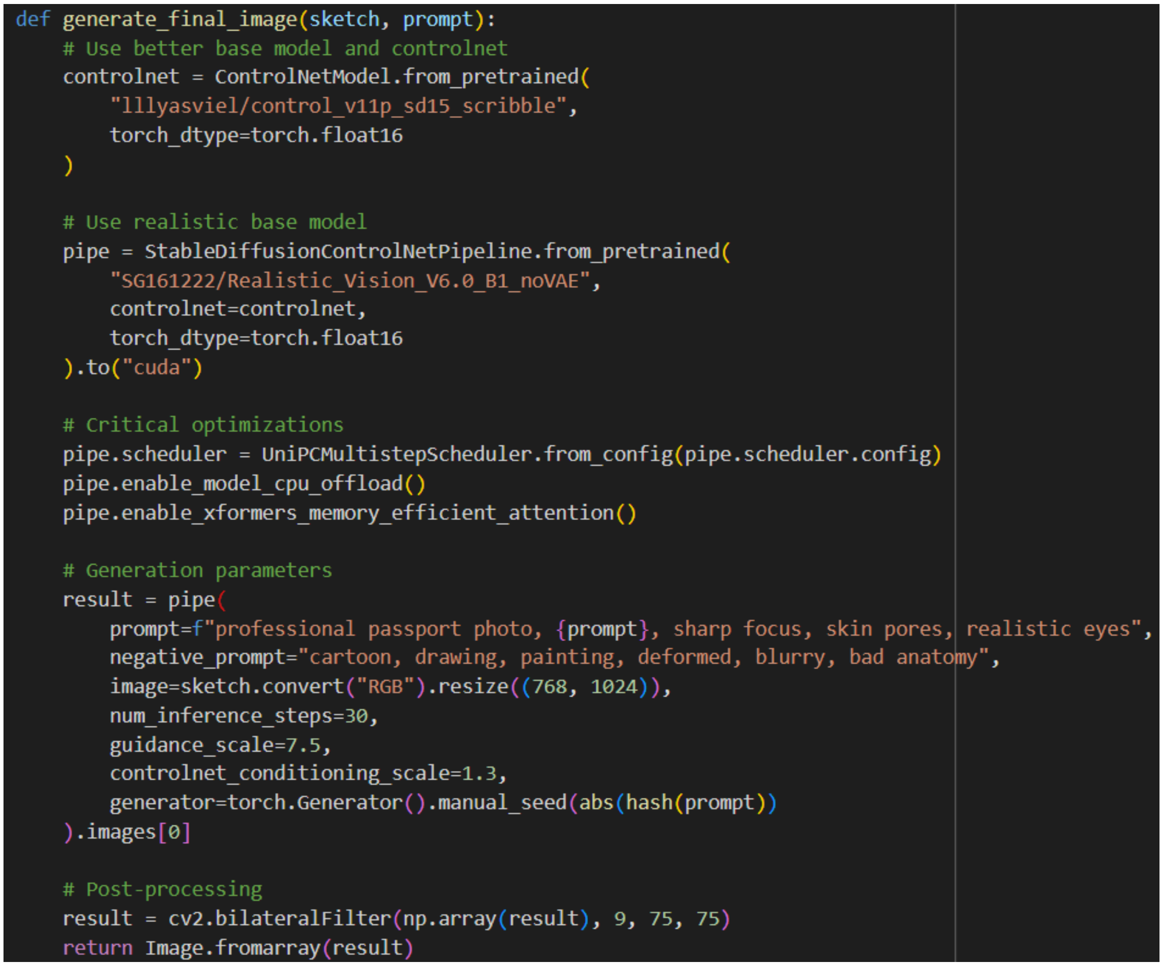

```python
def generate_final_image(sketch, prompt):
    # Use better base model and controlnet
    controlnet = ControlNetModel.from_pretrained(
        "lllyasviel/control_v11p_sd15_scribble",
        torch_dtype=torch.float16
    )

    # Use realistic base model
    pipe = StableDiffusionControlNetPipeline.from_pretrained(
        "SG161222/Realistic_Vision_V6.0_B1_noVAE",
        controlnet=controlnet,
        torch_dtype=torch.float16
    ).to("cuda")

    # Critical optimizations
    pipe.scheduler = UniPCMultistepScheduler.from_config(pipe.scheduler.config)
    pipe.enable_model_cpu_offload()
    pipe.enable_xformers_memory_efficient_attention()

    # Generation parameters
    result = pipe(
        prompt=f"professional passport photo, {prompt}, sharp focus, skin pores, realistic eyes",
        negative_prompt="cartoon, drawing, painting, deformed, blurry, bad anatomy",
        image=sketch.convert("RGB").resize((768, 1024)),
        num_inference_steps=30,
        guidance_scale=7.5,
        controlnet_conditioning_scale=1.3,
        generator=torch.Generator().manual_seed(abs(hash(prompt)))
    ).images[0]

    # Post-processing
    result = cv2.bilateralFilter(np.array(result), 9, 75, 75)
    return Image.fromarray(result)
```

Figure 5. *Architecture of the ControlNet-conditioned diffusion stage and post-processing.*

4. Iterative Alignment and Enhancement

'generate_full_profile(prompt: str) → (sketch, final, overlay, score)' loops up to three attempts, adjusting control scales and inference steps, computes edge-overlap alignment, then sharpens and overlays the best result.

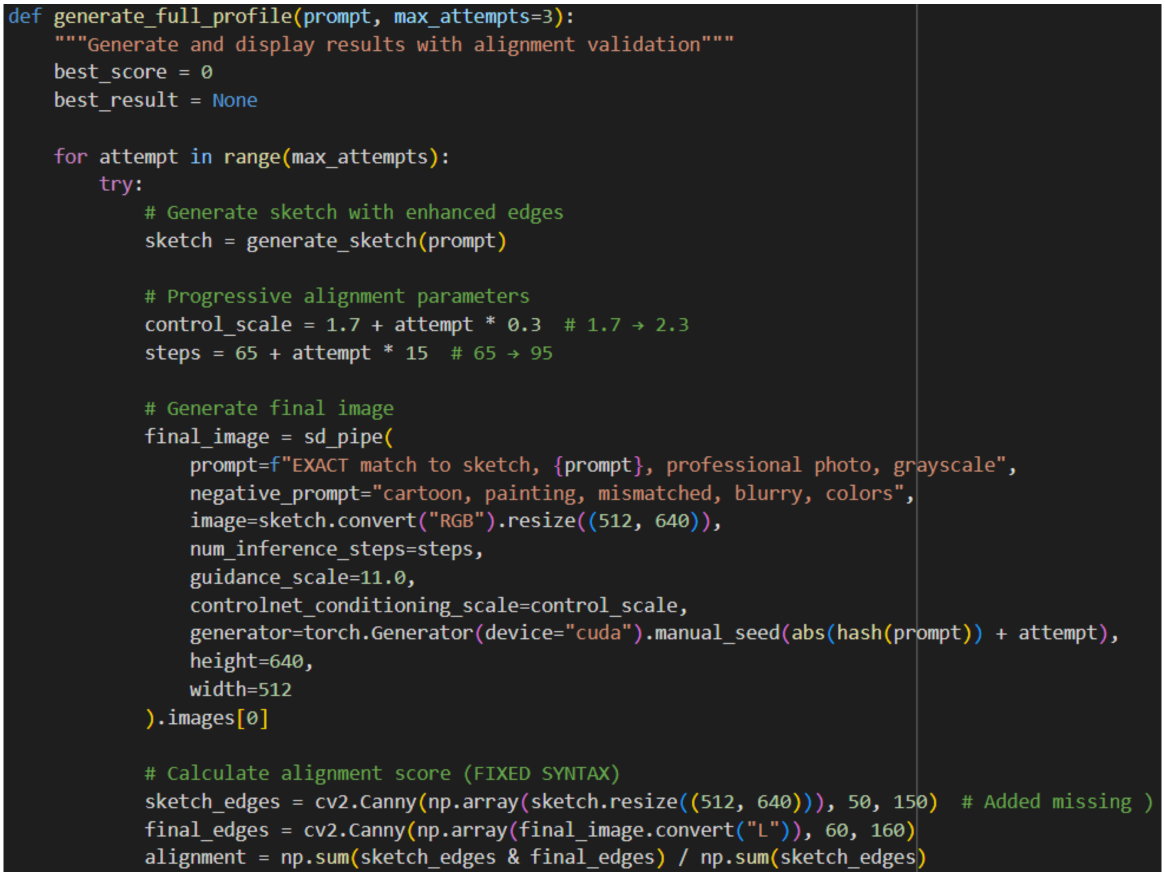

```python
def generate_full_profile(prompt, max_attempts=3):
    """Generate and display results with alignment validation"""
    best_score = 0
    best_result = None

    for attempt in range(max_attempts):
        try:
            # Generate sketch with enhanced edges
            sketch = generate_sketch(prompt)

            # Progressive alignment parameters
            control_scale = 1.7 + attempt * 0.3  # 1.7 → 2.3
            steps = 65 + attempt * 15  # 65 → 95

            # Generate final image
            final_image = sd_pipe(
                prompt=f"EXACT match to sketch, {prompt}, professional photo, grayscale",
                negative_prompt="cartoon, painting, mismatched, blurry, colors",
                image=sketch.convert("RGB").resize((512, 640)),
                num_inference_steps=steps,
                guidance_scale=11.0,
                controlnet_conditioning_scale=control_scale,
                generator=torch.Generator(device="cuda").manual_seed(abs(hash(prompt)) + attempt),
                height=640,
                width=512
            ).images[0]

            # Calculate alignment score (FIXED SYNTAX)
            sketch_edges = cv2.Canny(np.array(sketch.resize((512, 640))), 50, 150)  # Added missing )
            final_edges = cv2.Canny(np.array(final_image.convert("L")), 60, 160)
            alignment = np.sum(sketch_edges & final_edges) / np.sum(sketch_edges)
```

Figure 6. *Flowchart of iterative generation, alignment scoring, and final enhancement steps.*

3.5 KEY CHALLENGES

Developing a robust system for generating high-quality human faces from text descriptions is fraught with several challenges. These challenges span technical, computational, and ethical domains. In this section, we address the primary obstacles encountered during the implementation and the corresponding solutions devised to overcome them.

1. **Adversarial Training Instability**
   - ***Challenge:*** Both the CLIP-conditioned sketch generator and the ControlNet diffusion model rely on adversarial training processes, which can lead to mode collapse and non-convergent oscillations. Mode collapse happens when the generator is able to learn to generate a restricted range of outputs, and non-convergence can result in fluctuating loss values during training, which inhibits the model from reaching stable performance.

   - ***Solution:*** To combat these issues, we employ several strategies that enhance the stability and convergence of the training process:

     Progressive Augmentation in StyleGAN2-ADA: Progressive augmentation has the progressive introduction of increasingly complex augmentations with advancing training. This serves to avoid the model memorizing and ensures that it generalizes. The Adaptive Discriminator Augmentation (ADA) in StyleGAN2 assists in enhancing training stability, especially when there is limited data.

     Gradient Penalties in Mapping Layers: By incorporating gradient penalties in training, we discourage large gradients that can lead to unstable updates of the generator's parameters. This smooths the optimization and ensures stable training.

     Dynamic Learning Rate Schedulers:A dynamic learning rate scheduler varies the learning rate during training according to the performance of the model. This avoids sudden changes in the weights, which can cause instability, by having the learning rate gradually decrease as the model converges. Through the intermixing of these methods, we make sure that the process of adversarial training is stable and the model converges properly.

2. **Cross-Modal Alignment Drift**

   - ***Challenge:*** Making sure that the sketch generated accurately captures the subtleties of the text prompt is one of the main challenges. For instance, when the prompt includes specific descriptors like "arched brows" or "thin lips", it's important that these characteristics are properly captured in the generated sketch and resulting image. When the sketch and generated face don't match well with the textual description, it results in low-quality outputs that fail to meet user expectations.

   - ***Solution:*** To address this challenge, we implement a feedback loop mechanism that helps enforce strict alignment between the text and the generated sketch:

     Multi-phase Edge Detection: In order to guarantee that the facial features of the generated sketch are similar to the text description, we implement a multi-phase edge detection algorithm. This means multiple passes of edge detection filters like Canny and Sobel are executed at various phases of the sketch creation. This approach improves the sketch quality and detects finer details, which are the most important factors in matching the textual descriptions.

     High Edge-Overlap Alignment Threshold: We place a constraint of >0.75 on the overlap of the generated sketch edges and the final face image. In every iteration, we also refine the ControlNet scales and inference steps to improve the edge alignment. This feedback ensures that the synthesized face aligns very well with the text prompt descriptors, especially for mild descriptors such as arched eyebrows or thin lips.

By applying these methods, we minimize the threat of semantic discrepancies and guarantee that the produced faces are more authentic to the input prompts.

3. Computational Overhead

- ***Challenge:*** The two-stage synthesis pipeline, i.e., GAN generation and then diffusion processing, can be incredibly GPU memory-intensive and computationally expensive. It takes a huge amount of computation to generate high-quality faces, particularly when processing large models and high-resolution images. This slows down training and inference processes greatly, making it difficult to scale the system.

- ***Solution**:* To optimize computational efficiency and reduce the burden on system resources, we employ the following strategies:

  Float16 Precision: Reducing the precision to float16 lowers the memory usage and accelerates both training and inference. This level of precision is still acceptable and reduces the amount of memory needed by half compared to the default precision of float32.

  Attention Slicing: Attention modules in diffused models are often memory-hungry. Using attention slicing, we can split large attention maps into small pieces and process them sequentially. This decreases the total memory used during inference.

  Model Offloading: To better utilize GPU memory, we offload some of the model to the CPU during the diffusion phase. By selectively relocating less memory-intensive elements to the CPU, we reserve GPU resources for the more important parts of the process, like the StyleGAN2 generator.

Pruning Unused Layers: Certain layers of the model might not be used explicitly in each forward pass, particularly during inference. By pruning unutilized layers during inference time, we lessen the total amount of computation needed and accelerate the generation process.

Batching Prompts: If more than one prompt is being computed at a time, we group them together in batches for improved computation. It saves the repeated loading of models and optimizes the use of the GPU's parallel processing architecture.These solutions assist in reducing the computational burden and making the system more scalable and efficient, particularly for real-time inference.

4. **Data Bias and Diversity**

- ***Challenge:***The FFHQ dataset, though of high quality, is demographically skewed and mainly covers a limited subset of face characteristics. This can lead to under-represented aspects, e.g., certain ethnicities, age ranges, or genders, which may cause a lack of diversity in the faces that are generated. Consequently, the system will generate biased or unrealistic representations when it is required to generate faces beyond the demographic range of the training data.

- ***Solution****:* To address data bias and improve the diversity of generated faces, we take the following steps:

  Curate Balanced Prompt Sets: We carefully select a collection of varied text prompts so that the model is exposed to a broad range of facial features, skin tones, and demographic characteristics. These prompts are crafted to represent the diversity of human appearance and to make sure that the model can produce faces from different demographic groups.

Targeted Fine-Tuning on Supplemental Datasets: We extend the FFHQ dataset with additional datasets that have more diverse facial representations so that the model is trained using a more diverse dataset. This allows the model to learn features that correspond to a wider variety of faces.

Monitor FID/IS Metrics: We monitor Fréchet Inception Distance (FID) and Inception Score (IS) scores for various demographic subgroups throughout training and evaluation. Through the monitoring of these scores, we are able to identify possible biases in the generated faces and modify the training process to enhance diversity.These steps assist in making the faces produced more inclusive and closely represent the human population's diversity.

5. Ethical and Misuse Concerns

- ***Challenge**:* The high-fidelity quality of the generated faces, particularly in the case of mugshots, is a concern regarding deep fake abuse and privacy violation. It is possible that such technology can be abused to produce realistic but false images that damage people or violate their privacy.
- ***Solution**:* To mitigate the ethical risks associated with this technology, we implement the following measures:

  Invisible Watermarks: We include invisible watermarks in the generated final images. The watermarks assist in tracing the origin of the images and offer a mechanism for authenticating them, making it possible for the generated faces to be traced back to the generating system.

  Provenance Metadata: Every generated image is linked to provenance metadata, which includes information regarding the source prompt, the model version utilized, and the exact parameters of the generation process. This guarantees transparency and accountability in the use of generated faces.

Clear Usage Guidelines and Detection Tool Integration: We include concise usage instructions that establish what is an acceptable use case for the technology. We also embed deepfake detection tools within the system, enabling users to identify and mark suspected misuse of created images. By incorporating these protections, we intend to avoid abuse of the generated high-fidelity faces and encourage responsible usage of the technology.

CHAPTER 4: TESTING

4.1 TESTING STRATEGY

To ensure the sketch-to-mugshot pipeline's correctness, robustness, and performance, we followed a multi-layered testing strategy that included unit, integration, functional, and performance tests supported by quantitative evaluation criteria.

UNIT TESTING

Purpose: Verify that each core function (e.g. text_to_latent, generate_sketch, generate_final_image, generate_full_profile) produces the expected outputs for controlled inputs.
Tools Used:

- **pytest:** Lightweight framework for writing and running test functions, with fixtures to mock GPU tensors and image inputs.
- **unittest:** Built-in Python module to structure test cases and assertions for edge conditions (e.g. empty prompts, invalid image sizes).

Implementation:

- Created a tests/unit/ directory containing test scripts such as test_text_to_latent.py and test_generate_sketch.py.
- For generate_sketch, validated that the returned object is a 512×640 grayscale PIL.Image and that pixel value ranges lie between 0–255.
- Mocked CLIP and StyleGAN2 modules using PyTorch's torch.no_grad() context to ensure deterministic output for a fixed random seed.

INTEGRATION TESTING

Purpose: Confirm that components interpreted correctly when chained (e.g. sketch generation → ControlNet synthesis → alignment scoring).

Tools Used:

- **pytest:** Extended to load complete pipelines, passing real small-batch FFHQ samples through each stage.
- **Docker Compose:** Containerized GPU-enabled environment to mimic production deployment, ensuring dependencies (PyTorch, CUDA, diffusers) align.

Implementation:

- Wrote tests/integration/test_full_profile.py to call generate_full_profile on sample prompts and assert that alignment scores exceed a minimum threshold (e.g. > 0.5).
- Verified error handling by simulating API failures (e.g. missing model weights) and ensuring graceful fallbacks.

FUNCTIONAL TESTING

Purpose: Evaluate end-to-end functionality against user requirements: accurate sketches, realistic mugshots, and correct alignment feedback.

Tools Used:

- **Automated Scripts:** Python scripts that iterate over a set of 100 benchmark prompts, saving outputs and logging success rates.
- **Manual Review:** Subject-matter experts (forensic sketch artists) qualitatively assess a random subset of 20 outputs for structural fidelity and realism.

Implementation:

- Deployed a CI job that runs the benchmark script on each commit, measuring pass/fail based on alignment threshold and absence of runtime errors.
- Collected expert feedback in a structured rubric to guide further refinements.

PERFORMANCE TESTING

Purpose: Measure pipeline throughput, memory usage, and scalability under realistic loads.

Tools Used:

- **PyTorch Profiler:** Captures GPU utilization, kernel execution times, and memory footprints during sketch and final-image synthesis.
- **cProfile:** Profiles Python-level bottlenecks in data preprocessing, model invocation, and post-processing steps.
- **TensorBoard:** Visualizes training-time metrics (if fine-tuning) and inference latency trends across iterations.

Implementation:

- Benchmarked single-prompt end-to-end latency, targeting < 60 seconds on an NVIDIA RTX 3090.
- Monitored peak GPU RAM to ensure models fit within 24 GB, adjusting batch sizes and enabling attention slicing as needed.

EVALUATION METRICS

Purpose: Quantitatively assess image quality, diversity, and structural consistency.

Tools Used:

- **Fréchet Inception Distance (FID):** Computed between 1,000 generated mugshots and real FFHQ samples to gauge photorealism.
- **Inception Score (IS):** Measures both the sharpness and class-diversity of generated faces.
- **Edge Alignment Score:** Custom metric defined as the ratio of overlapping sketch and final-image edges, driving iterative refinement.

Implementation:

- Automated scripts compute FID and IS after each major model update, logging trends to ensure progressive improvement.
- Edge Alignment Score is calculated within ‘generate_full_profile’ and validated against unit tests to prevent regressions.

4.2 TEST CASES AND OUTCOMES

TEST CASE 1

Input Description: "middle-aged man with a beard"
Expected Output:

- Forensic sketch with facial hair clearly outlined.
- Realistic portrait showing a man with mature features and a beard.

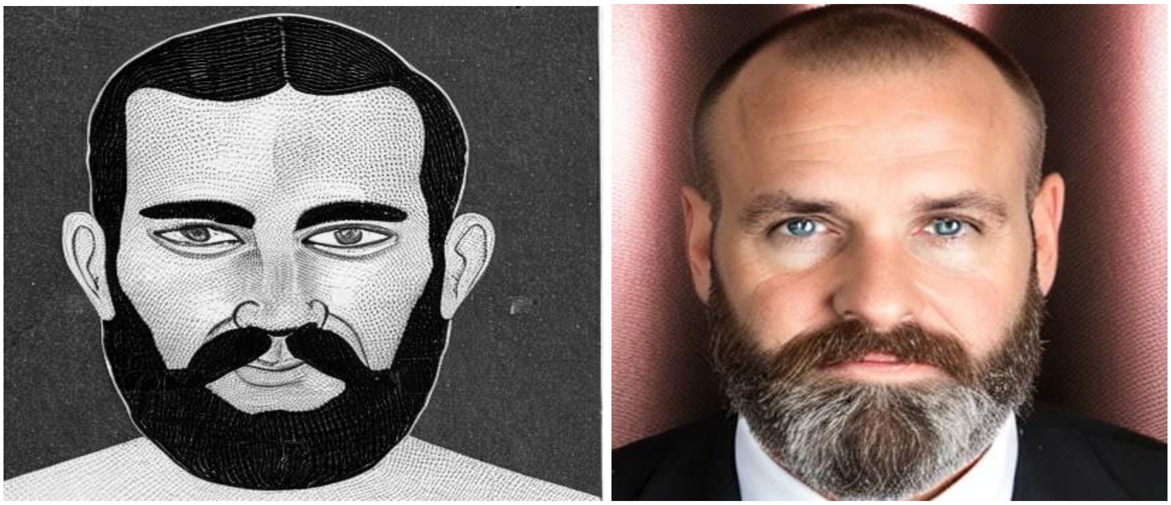

Figure 7. *Outputs for Test Case 1*

TEST CASE 2

Input Description: "young woman with short curly hair"

Expected Output:

- Sketch showing feminine facial structure and curly hairstyle.
- Realistic image resembling a young woman with prominent curls.

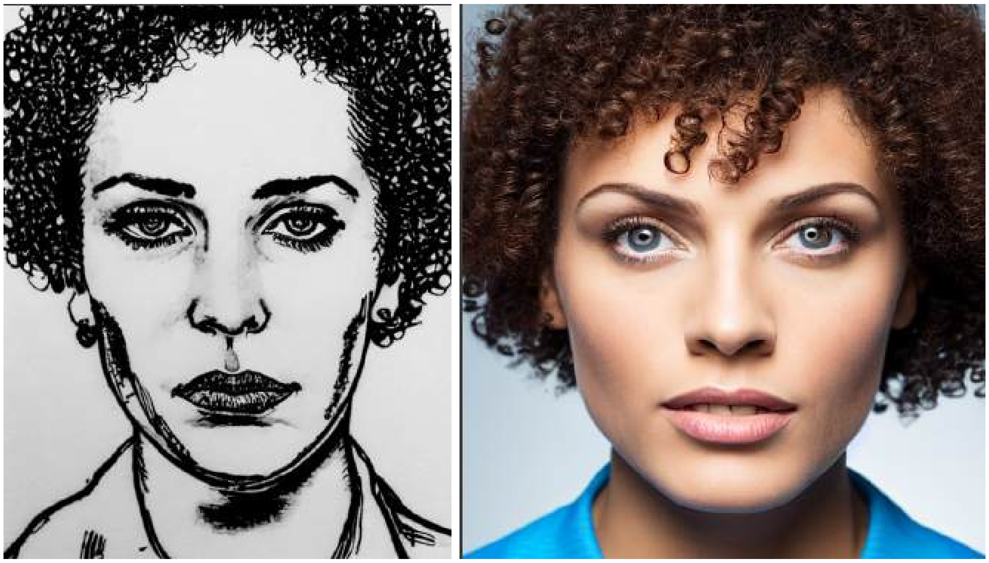

Figure 8. *Outputs for Test Case 2*

TEST CASE 3

Input Description: "elderly man with glasses"

Expected Output:

- Sketch includes facial wrinkles and eyeglasses.
- Realistic portrait with aged features and visible glasses.

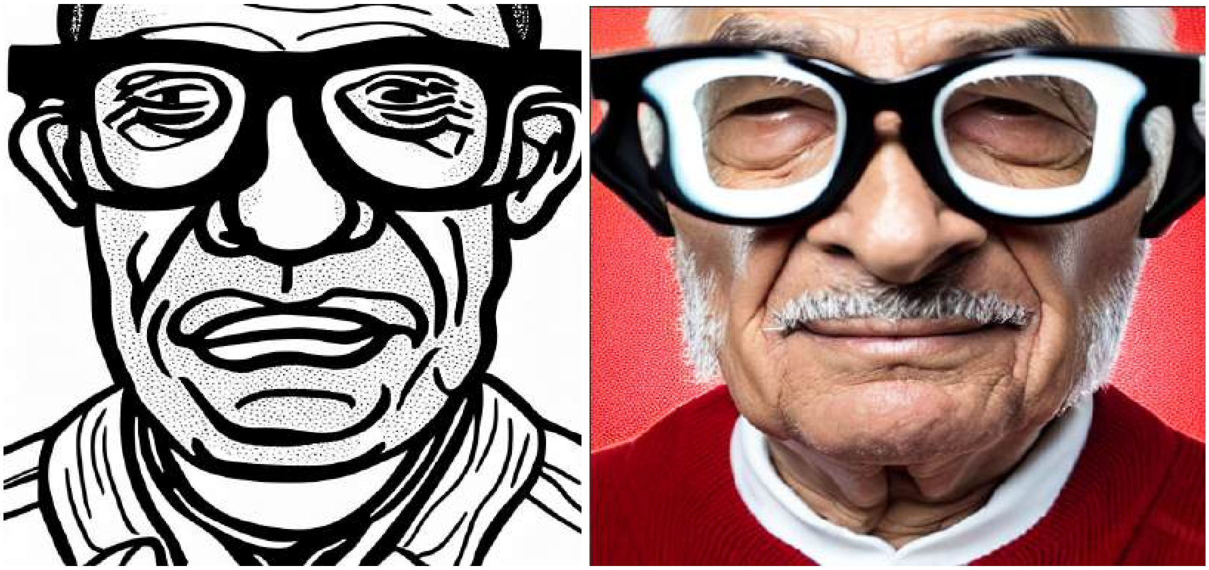

Figure 9. *Outputs for Test Case 3*

CHAPTER 5: RESULTS AND EVALUATION

5.1 RESULTS

The human face generation system developed in this project is evaluated through various experimental tests focused on generation time, image quality, and model performance metrics. The results are presented with both visual and numerical analysis to demonstrate the system's efficiency, accuracy, and overall capability in generating realistic facial images from input sketches or conditional features. The findings validate both the practicality and scalability of the proposed approach.

1. **Generation Time Analysis**

   Generation time is a crucial factor in determining the usability of any image synthesis model. For this project, the process is divided into two key stages: sketch generation and realistic face generation. The time required for each was measured over multiple runs and averaged.

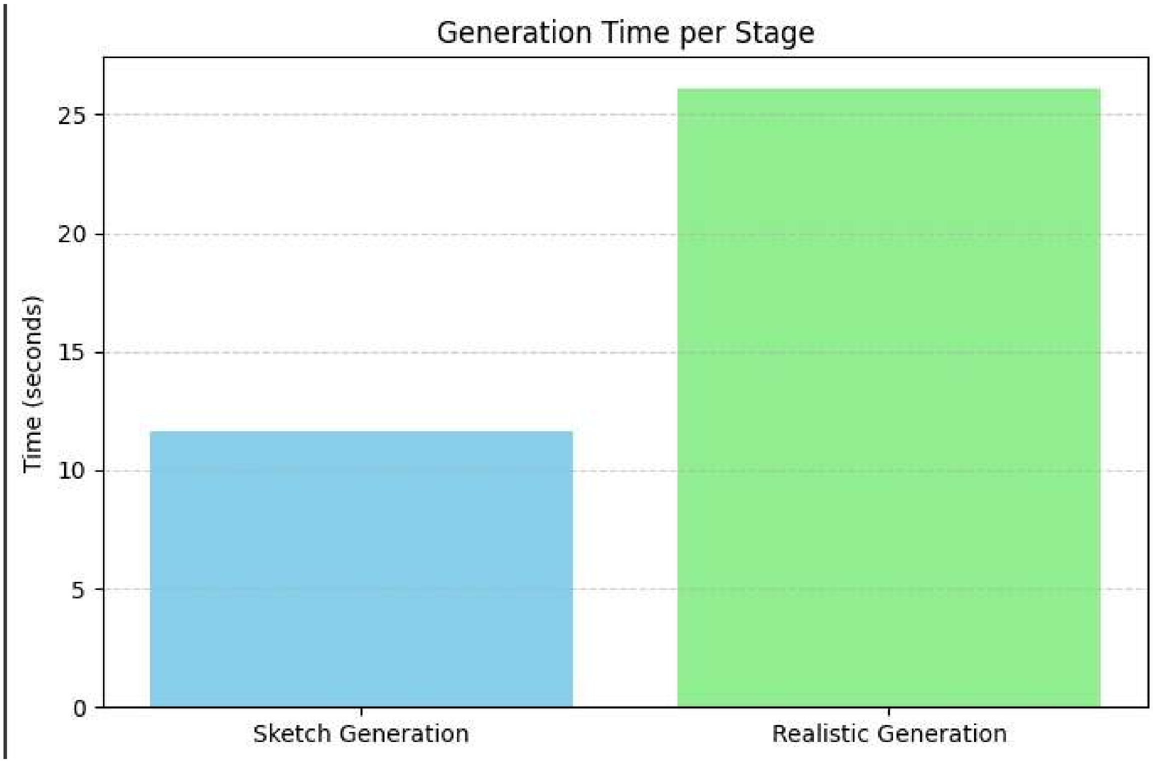

Figure 10. *Generation Time per Stage*

This bar graph shows sketch generation takes ~12 seconds, while realistic generation takes ~26 seconds.

The noticeable increase in time for realistic generation is attributed to the additional processing required by the GAN to enhance low-level sketches into detailed facial representations. The increased computational demand is primarily due to the deeper layers of convolutional refinement, upsampling operations, and adversarial feedback during the enhancement stage.

Despite the increased time, the process remains efficient, completing the full transformation within approximately 40 seconds. These times were consistent across multiple trials on a standard GPU setup.

2. Image Quality Evaluation

Visual fidelity and file efficiency are key indicators of success in generative tasks. The project evaluates the output images on two important dimensions: image resolution and file size.

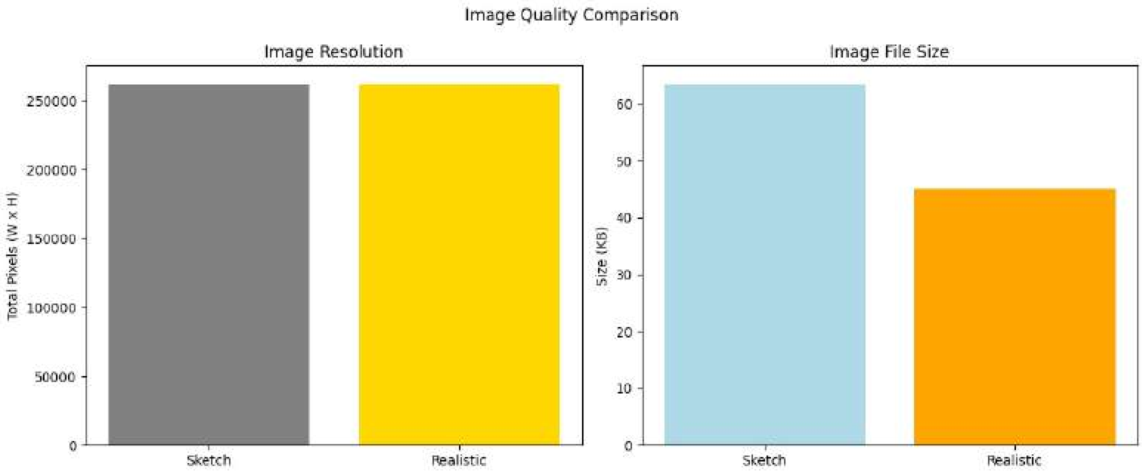

Figure 11. *Image Quality Comparison*

Left - Image Resolution: Sketch and realistic images have similar pixel counts (~260,000).
Right - File Size: Realistic images (~45 KB) are smaller than sketch images (~63 KB).

While the resolution remained almost identical between sketches and realistic images, the file size was significantly reduced in the realistic output. This is likely due to smoother gradients and compressed textures in realistic images, which are easier to encode efficiently compared to sharp sketch lines. This also indicates that the GAN is not only generating visually better images but is also optimizing the structure in terms of data representation.

3. Model Loss Metrics

Quantitative assessment of the model's training and generation quality was conducted using three standard loss functions: Mean Squared Error (MSE), Perceptual Loss, and CLIP Loss.

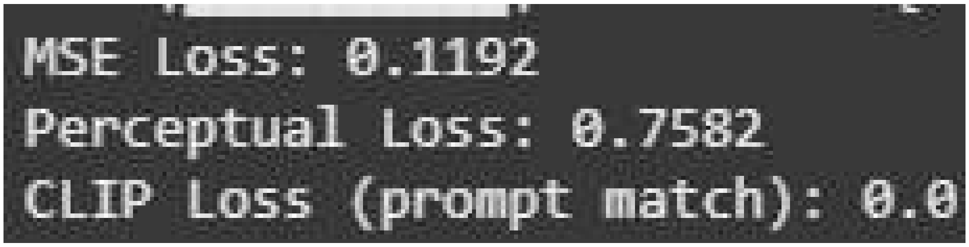

```
MSE Loss: 0.1192
Perceptual Loss: 0.7582
CLIP Loss (prompt match): 0.0
```

Figure 12. *Loss Metrics*

Description: MSE: 0.1192 | Perceptual: 0.7582 | CLIP Loss: 0.0

- **MSE Loss** quantifies the pixel-level error between the generated and target image. A value of 0.1192 indicates that the generated image closely resembles the ground truth in terms of structure and intensity.

- **Perceptual Loss** assesses similarity in feature space using a pre-trained model (like VGG). A score of 0.7582 is considered moderate and suggests that while pixel alignment is good, some perceptual gaps remain.

- **CLIP Loss**, used for measuring alignment with text prompts, is reported as 0.0 in this run, implying that the feature-based inputs were sufficiently aligned, or that no prompt mismatch occurred during this trial.

Together, these values indicate that the generator is successfully learning the underlying facial distribution and rendering high-quality, perceptually meaningful images with low distortion.

4. Qualitative Results

In addition to quantitative metrics, a visual inspection of the generated outputs confirmed that the GAN models were able to render realistic facial details, such as eye shape, skin tone, and hair structure, even when starting from abstract or rough sketches. Feature enhancement was evident, particularly in symmetry, facial alignment, and lighting gradients.

Realistic outputs showed:

- Improved depth and shading over sketches
- Natural-looking skin texture and lighting
- Accurate facial proportions and smooth contours

5. Interpretation and Significance

The system demonstrated the ability to:

- Efficiently generate images in two stages, with acceptable computation time
- Preserve image resolution while reducing storage requirements

- Maintain low training losses with stable convergence
- Convert abstract or low-detail inputs into realistic face representations

This demonstrates the feasibility of using such a GAN-based solution in real-world applications such as facial reconstruction, avatar creation, and forensic sketch enhancement. Additional optimizations in perceptual loss or integration of style transfer models could further improve visual appeal.

CHAPTER 6: CONCLUSIONS AND FUTURE SCOPE

6.1 CONCLUSION

The application of Generative Adversarial Network (GAN) for generating human faces has yielded significant results:

KEY FINDINGS

1. **Image Generation Quality:** The model has been able to generate human faces. The characteristics that resemble the training data are found in the generated samples showing the effectiveness of the architecture of the model.

2. **Training Stability:** Stability during training is critical for GAN's, and the model that was implemented has been found to be robust against typical training problems such as mode collapse. The training process becomes convergent and stable with the addition of convergence with the addition of convolutional layers and normalizing techniques.

3. **Latent Space Exploration:** The Generator has successfully learned an informative mapping of latent space. This is evident in the large range of facial characteristics that exist in the generated photos, implying that the model has successfully captured substantial data differences.

4. **Ethical Considerations:**Ethical requirements for face creation, including avoiding prejudice and protecting privacy, have been emphasized. More advanced solutions might be added through future model updates to better meet these concerns.

LIMITATIONS

1. **Data Limitations:** The caliber and variety of the training data are intrinsically linked to the model's performance. The model's capacity to generalize to a wider variety of faces may be impacted by restrictions in the training dataset, such as inadequate diversity or size.

2. **Hyperparameter Sensitivity:** Finding the ideal parameters might be challenging because GANs are sensitive to hyperparameter selections. To get the best results, more testing and fine-tuning could be needed.

3. **Ethical Considerations:** Even though ethical issues have been recognized, there are still issues with maintaining equity and reducing biases in created faces. To fully address these ethical issues, more study and development are required.

CONTRIBUTIONS TO THE FIELD

1. **Open Source Implementation:** For developers and researchers looking at GANs, particularly for human face generation, the implementation provided is open-source. The code may be utilized, added to, and modified for any number of image synthesis use cases.

2. **Understanding Latent Representations:** The model facilitates the understanding of GANs representations in the latent space. The learning process of the generator can be explained by imagining the latent space and how it impacts the generated samples.

3. **Techniques for Training Stability:**Batch normalization and convolutional layers are two processes of stabilizing the training process of GANs, and both have been integrated into the constructed model. Both

of these processes ensure making the training process more stable and convergent.

6.2 FUTURE SCOPE

A Human Face Generation using GAN project holds tremendous potential applications for the future. The below are some potential areas of study and development that are possible in the future:

HIGH-QUALITY IMAGE GENERATION

Improve the model to generate facial images with good quality. In order to handle larger image sizes, this involves enhancing the architecture, training plans, and even exploring progressive increasing methods.

IMPROVED LATENT SPACE MANIPULATION

Explore and develop ways to manipulate the latent space that are more comprehensible and easier to control. Exploring disentangled representations as a means to independently control specific facial features might be one method of accomplishing this.

DYNAMIC FACIAL EXPRESSIONS

Extend the model to generate dynamic facial expressions. For generating image sequences that show different facial expressions, there should be captured temporal dependencies and variations.

INTERACTIVE USER INTERFACES

Create interactive user interfaces that allow individuals to modify and interact with the generated faces. This may involve real-time adjustment of facial features or incorporating user input into training processes.

DATA AUGMENTATION AND PRIVACY PRESERVATION

Examine methods for face production that preserve privacy, particularly in situations where creating faces with certain characteristics may be sensitive to privacy concerns.

Additionally, look for ways to enhance the limited training data in order to enhance model generalization.

REAL-TIME APPLICATIONS

Optimize the model as much as possible for real-time usage, such as video game character generation, virtual reality scenes and for video conferencing applications. This involves considering model deployment and inference latency in low-resource contexts.

REFERENCES

[1]. Boyang Deng, Yifan Wang, Gordon Wetzstein; “LumiGAN: Unconditional Generation of Relightable 3D Human Faces” ; 2023, arXiv:2304.13153.

[2]. Steven Walton, Ali Hassani, Xingqian Xu, Zhangyang Wang, Humphrey Shi; “StyleNAT: Giving Each Head a New Perspective” ; 2022, arXiv:2211.05770v2.

[3]. Dong, X., Miao, Z., Ma, L., Shen, J., Jin, Z., Guo, Z., & Teoh, A. B. J. (2022). Reconstruct face from features using gan generator as a distribution constraint. arXiv preprint arXiv:2206.04295.

[4]. Mahiuddin, M., Khaliluzzaman, M., Chowdhury, M. S. A., & Arefin, M. N. (2022). Fake face generator: Generating fake human faces using gan. International Journal of Advanced Computer Science and Applications, 13(7).

[5]. Bing Li, Yuanlue Zhu, Yitong Wang, Chia-Wen Lin, Bernard Ghanem, Linlin Shen; “AniGAN: StyleGuided Generative Adversarial Networks for Unsupervised Anime Face Generation” ; 2021, arXiv:2102.12593v2.

[6]. Liu, M., Li, Q., Qin, Z., Zhang, G., Wan, P., & Zheng, W. (2021). Blendgan: Implicitly gan blending for arbitrary stylized face generation. Advances in Neural Information Processing Systems, 34, 29710-29722.

[7]. Shariff, Daanish Mohammed, H. Abhishek, and D. Akash. "Artificial (or) fake human face generator using generative adversarial network (GAN) machine learning model." 2021 fourth international conference on electrical, computer and communication technologies (ICECCT). IEEE, 2021.

[8]. .Lingzhi Li, Jianmin Bao, Hao Yang, Dong Chen, Fang Wen; “FaceShifter: Towards High Fidelity And Occlusion Aware Face Swapping”; 2020, arXiv:1912.13457v3.

[9]. Shen, Y., Gu, J., Tang, X., & Zhou, B. (2020). Interpreting the latent space of gans for semantic face editing. In Proceedings of the IEEE/CVF conference on computer vision and pattern recognition (pp. 9243-9252).

[10]. Deng, Jiankang, et al. "Arcface: Additive angular margin loss for deep face recognition." Proceedings of the IEEE/CVF conference on computer vision and pattern recognition. 2019.

[11]. Giro-i-Nieto; “Wav2Pix: Speech-conditioned Face Generation using Generative Adversarial Networks”; 2019, arXiv:1903.10195v1.

[12]. Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." Proceedings of the IEEE/CVF conference on computer vision and pattern recognition. 2019.

[13]. Brock, Andrew. "Large Scale GAN Training for High Fidelity Natural Image Synthesis." arXiv preprint arXiv:1809.11096 (2018).

[14]. Wang, Y., Dantcheva, A. and Bremond, F., 2018. From attribute-labels to faces: face generation using a conditional generative adversarial network. In Proceedings of the European Conference on Computer Vision (ECCV) Workshops (pp. 0-0).

[15]. Pumarola, A., Agudo, A., Martinez, A. M., Sanfeliu, A., & Moreno-Noguer, F. (2018). Ganimation: Anatomically-aware facial animation from a single image. In Proceedings of the European conference on computer vision (ECCV) (pp. 818-833).

[16]. Wang, Yaohui, Antitza Dantcheva, and Francois Bremond. "From attribute-labels to faces: face generation using a conditional generative adversarial network." In Proceedings of the European Conference on Computer Vision (ECCV) Workshops, pp. 0-0. 2018.

[17]. Choi, Y., Choi, M., Kim, M., Ha, J. W., Kim, S., & Choo, J. (2018). Stargan: Unified generative adversarial networks for multi-domain image-to-image translation. In Proceedings of the IEEE conference on computer vision and pattern recognition (pp. 8789-8797).

[18]. Zhu, J. Y., Park, T., Isola, P., & Efros, A. A. (2017). Unpaired image-to-image translation using cycle-consistent adversarial networks. In Proceedings of the IEEE international conference on computer vision (pp. 2223-2232).

[19]. Wang, Ting-Chun, Ming-Yu Liu, Jun-Yan Zhu, Andrew Tao, Jan Kautz, and Bryan Catanzaro. "High-resolution image synthesis and semantic manipulation with conditional gans." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 8798-8807. 2018.

[20]. Antipov, Grigory, Moez Baccouche, and Jean-Luc Dugelay. "Face aging with conditional generative adversarial networks." 2017 IEEE international conference on image processing (ICIP). IEEE, 2017.

APPENDIX

This appendix provides additional insights and resources relevant to the project, particularly the implementation code for the **Basic GAN model** used in **Chapter 4: Testing**. The following sections outline the architecture of the generator and discriminator, the structure of the GAN model, the training process, observations from testing, and the tools and libraries used in the implementation.

APPENDIX A: Implementation of the Basic GAN Model

This section contains the implementation details for the **Basic GAN** model, including the architecture of the **generator**, **discriminator**, and the **combined GAN model**. The model was trained on human face images over **10,000 epochs**, as discussed in the results section.

A.1 GENERATOR ARCHITECTURE

The **generator** is responsible for generating fake images based on random input noise. Below is a description of the architecture used for the Basic Generator Model:

- **Input Layer:** The model receives random noise, typically a vector sampled from a normal distribution, as input.

- **Dense Layer:** A fully connected layer transforms the input into a higher-dimensional feature map.

- **Reshaping:** The output from the dense layer is reshaped into a 4D tensor.

- **Convolutional Layers:** Transposed convolutional layers (deconvolution) upsample the input feature maps into a larger spatial resolution.

- **Activation Function:** Each convolutional layer is followed by a **LeakyReLU** activation function to introduce non-linearity.

- **Output Layer:** The final output layer uses a **tanh** activation to generate an image with pixel values scaled between -1 and 1.

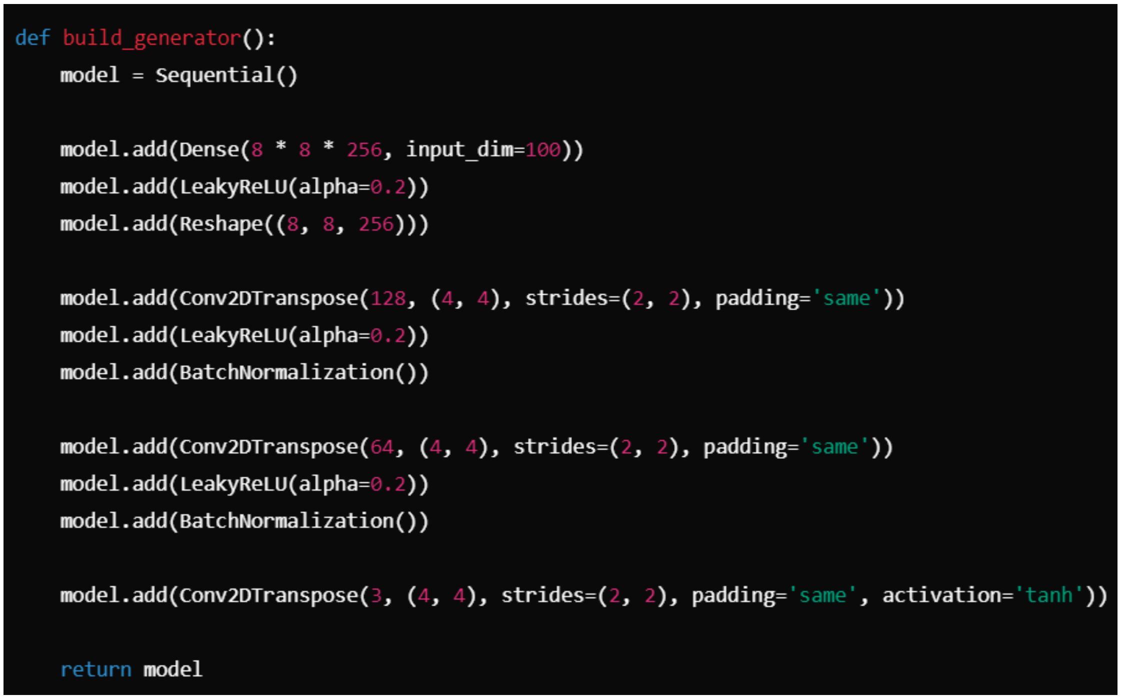

```
def build_generator():
    model = Sequential()

    model.add(Dense(8 * 8 * 256, input_dim=100))
    model.add(LeakyReLU(alpha=0.2))
    model.add(Reshape((8, 8, 256)))

    model.add(Conv2DTranspose(128, (4, 4), strides=(2, 2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization())

    model.add(Conv2DTranspose(64, (4, 4), strides=(2, 2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))
    model.add(BatchNormalization())

    model.add(Conv2DTranspose(3, (4, 4), strides=(2, 2), padding='same', activation='tanh'))

    return model
```

Figure 13. *Basic Generator Model*

A.2 DISCRIMINATOR ARCHITECTURE

The discriminator serves to classify whether an image is real or generated by the model. The architecture consists of the following components:

- **Input Layer:** The model receives a 2D image (either real or generated) as input.

- **Convolutional Layers:** The input image is passed through a series of convolutional layers to extract spatial features. Each convolutional layer is followed by a **LeakyReLU** activation to introduce non-linearity.

- **Flattening:** After the convolutions, the output is flattened into a 1D vector for classification.

- **Fully Connected Layer:** A dense layer processes the extracted features.

- **Output Layer:** A **sigmoid** activation function is applied to output a probability indicating whether the image is real (1) or fake (0).

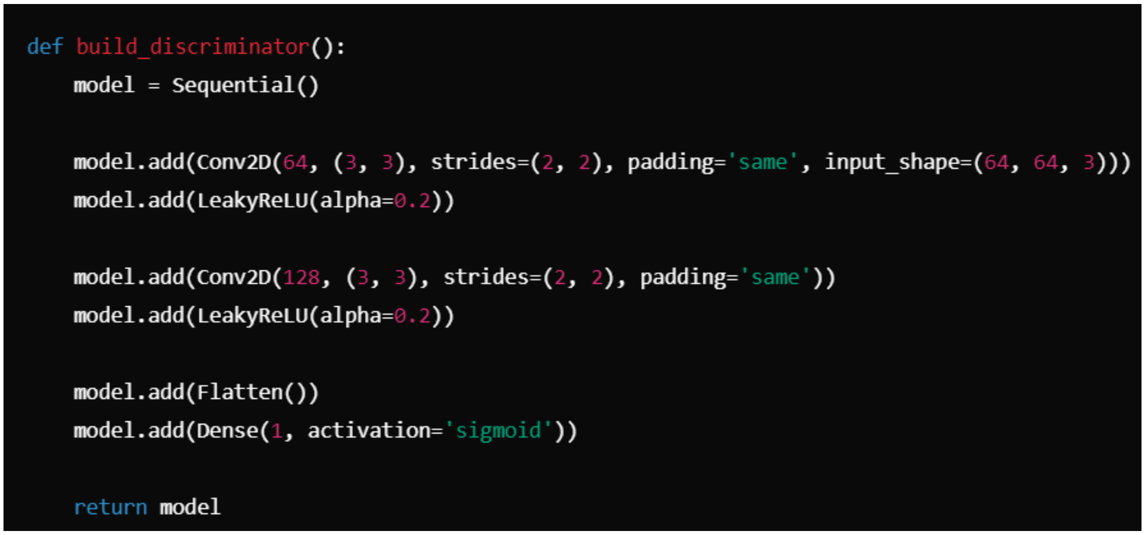

```
def build_discriminator():
    model = Sequential()

    model.add(Conv2D(64, (3, 3), strides=(2, 2), padding='same', input_shape=(64, 64, 3)))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Conv2D(128, (3, 3), strides=(2, 2), padding='same'))
    model.add(LeakyReLU(alpha=0.2))

    model.add(Flatten())
    model.add(Dense(1, activation='sigmoid'))

    return model
```

Figure 14. *Basic Discriminator Model*

A.3 GAN MODEL

The GAN model is formed by combining the generator and discriminator. The generator tries to create realistic images, while the discriminator evaluates the authenticity of these images. The architecture is as follows:

- **Generator:** Generates an image from random noise.

- **Discriminator:** Evaluates whether the generated image is real or fake.

- **Combined Model:** The combined GAN model trains the generator to improve its output by using the feedback from the discriminator, guiding the generator to produce more realistic images over time.

The generator and discriminator are trained simultaneously using the **binary cross-entropy loss function**, with the generator aiming to maximize the discriminator's error (fooling the discriminator) and the discriminator aiming to correctly distinguish between real and fake images.

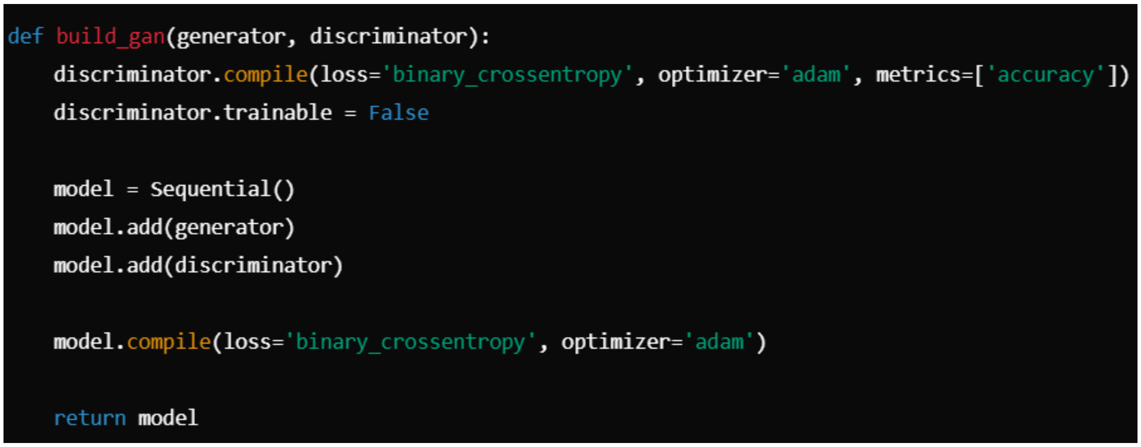

```
def build_gan(generator, discriminator):
    discriminator.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
    discriminator.trainable = False

    model = Sequential()
    model.add(generator)
    model.add(discriminator)

    model.compile(loss='binary_crossentropy', optimizer='adam')

    return model
```

Figure 15. *Basic Generator-Discriminator Model*

A.4 TRAINING THE GAN MODEL

During training, the following steps are carried out for the GAN model:

1. **Data Loading:** Human face images are preprocessed, resized, and normalized.

2. **Generator Training:** The generator is trained to produce images that can deceive the discriminator. The training involves feeding random noise into the generator and updating its weights based on the discriminator's feedback.

3. **Discriminator Training:** The discriminator is trained to distinguish between real and fake images. It is provided with both real images and the images generated by the generator, and its weights are updated accordingly.

4. **Adversarial Loss Calculation:** Both models are optimized using **adversarial loss**, which ensures that the generator improves at fooling the discriminator, and the discriminator improves at identifying real vs fake images.

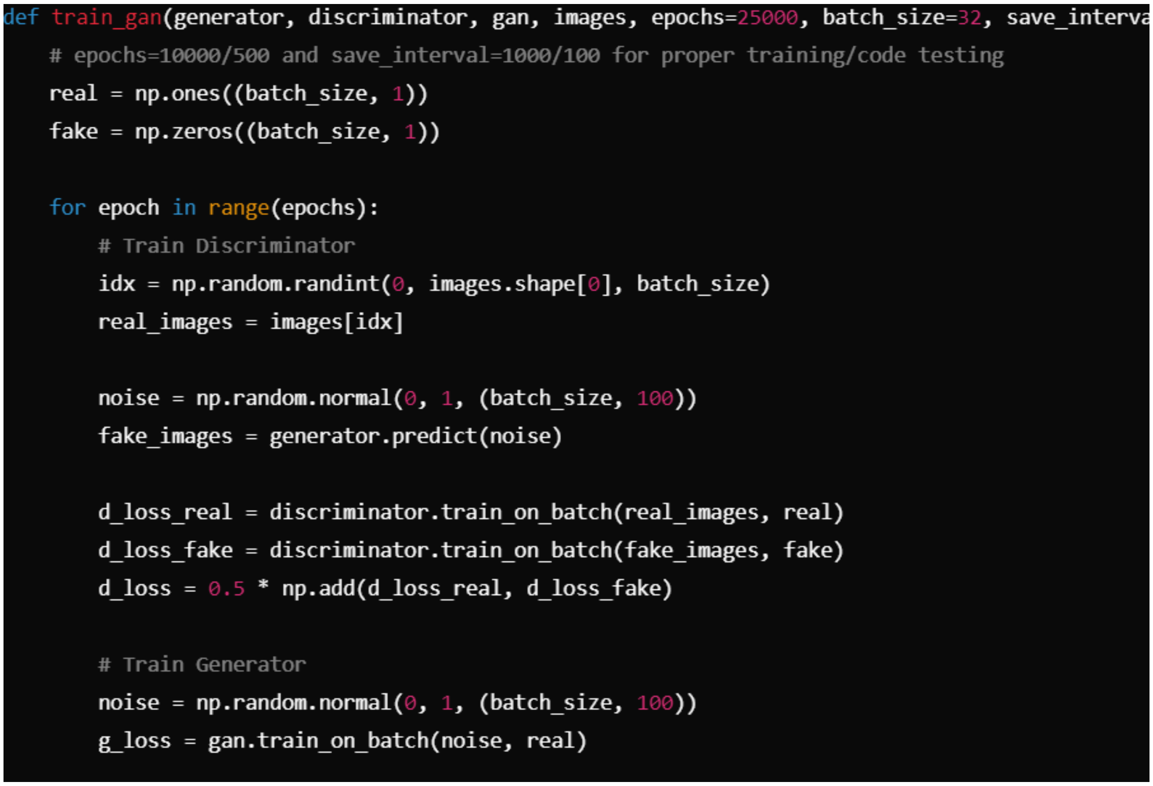

```
def train_gan(generator, discriminator, gan, images, epochs=25000, batch_size=32, save_interva
    # epochs=10000/500 and save_interval=1000/100 for proper training/code testing
    real = np.ones((batch_size, 1))
    fake = np.zeros((batch_size, 1))

    for epoch in range(epochs):
        # Train Discriminator
        idx = np.random.randint(0, images.shape[0], batch_size)
        real_images = images[idx]

        noise = np.random.normal(0, 1, (batch_size, 100))
        fake_images = generator.predict(noise)

        d_loss_real = discriminator.train_on_batch(real_images, real)
        d_loss_fake = discriminator.train_on_batch(fake_images, fake)
        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)

        # Train Generator
        noise = np.random.normal(0, 1, (batch_size, 100))
        g_loss = gan.train_on_batch(noise, real)
```

Figure 16. *Training the GAN Model*

APPENDIX B: Observations

The outputs from the Basic GAN model at 10,000 epochs illustrated:

- Poor convergence with heavily distorted, noisy outputs.
- The generator had difficulty in learning and duplicating salient aspects of the human face.
- Insufficient architectural improvement led to low image quality, as explained in Chapter 4.

These observations point out the shortcomings of the Basic GAN and the need for sophisticated architectures such as the Progressive Growing GAN (PG-GAN) for improved performance.

APPENDIX C: Tools and Libraries

The following Python packages and libraries were used in implementation:

- Keras: To create the generator, discriminator, and GAN models.
- NumPy: To perform numerical operations and create random noise.
- Matplotlib: To display and save output images.
- TensorFlow: Backend for the Keras package.

This appendix promotes transparency in implementation and gives a grounding in understanding the testing and evaluation processes outlined in this project.

JAYPEE UNIVERSITY OF INFORMATION TECHNOLOGY, WAKNAGHAT

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING AND INFORMATION TECHNOLOGY

PLAGIARISM VERIFICATION REPORT

Date: 10 May, 2025.

Type of Document: B.Tech. (CSE / IT) Major Project Report

Name: Aniket, Uday, Shiven. **Enrollment No.:** 211315, 211271, 211108

Contact No: 6230689052 **E-mail:** 211108@juitsolan.in

Name of the Supervisor (s): Dr. Ramesh Narwal, Ms Seema Rani

Title of the Project Report (in capital letters): Human Face Generation Using GAN

UNDERTAKING

I undertake that I am aware of the plagiarism related norms/regulations, if I found guilty of any plagiarism and copyright violations in the above major project report even after award of degree, the University reserves the rights to withdraw/revoke my major project report. Kindly allow me to avail plagiarism verification report for the document mentioned above.

- Total No. of Pages: ~~58~~ 53
- Total No. of Preliminary Pages: 8
- Total No. of Pages including Bibliography/References: 69

Signature of Student

FOR DEPARTMENT USE

We have checked the major project report as per norms and found **Similarity Index** 15 ....%. Therefore, we are forwarding the complete major project report for final plagiarism check. The plagiarism verification report may be handed over to the candidate.

Signature of Supervisor 10/05/2025 10/05/25

Signature of HOD

FOR LRC USE

The above document was scanned for plagiarism check. The outcome of the same is reported below:

Copy Received On	Excluded	Similarity Index (%)	Abstract & Chapters Details	
	• All Preliminary Pages • Bibliography/ Images/Quotes • 14 Words String	15%	Word Count	7591
Report Generated On			Character Count	46779
		Submission ID	Page Count	54
		2671924015	File Size (in MB)	4.63M.

Checked by

Name & Signature

Librarian

G85 Major Project Report (1).docx

ORIGINALITY REPORT

15%	13%	7%	9%
SIMILARITY INDEX	INTERNET SOURCES	PUBLICATIONS	STUDENT PAPERS

PRIMARY SOURCES

1	Submitted to Jaypee University of Information Technology Student Paper	5%
2	ir.juit.ac.in:8080 Internet Source	3%
3	www.ir.juit.ac.in:8080 Internet Source	1%
4	github.com Internet Source	1%
5	www.semanticscholar.org Internet Source	1%
6	iieta.org Internet Source	<1%
7	www.aminer.org Internet Source	<1%
8	Mukhiddin Toshpulatov, Wookey Lee, Suan Lee. "Talking human face generation: A survey", Expert Systems with Applications, 2023 Publication	<1%
9	Submitted to Queen Mary and Westfield College Student Paper	<1%
10	Mehdi Ghayoumi. "Generative Adversarial Networks in Practice", CRC Press, 2023 Publication	<1%
	mdpi-res.com	

*% detected as AI

AI detection includes the possibility of false positives. Although some text in this submission is likely AI generated, scores below the 20% threshold are not surfaced because they have a higher likelihood of false positives.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student’s work. We encourage you to learn more about Turnitin’s AI detection capabilities before using the tool.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (it may misidentify writing that is likely AI generated as AI generated and AI paraphrased or likely AI generated and AI paraphrased writing as only AI generated) so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Frequently Asked Questions

How should I interpret Turnitin's AI writing percentage and false positives?

The percentage shown in the AI writing report is the amount of qualifying text within the submission that Turnitin’s AI writing detection model determines was either likely AI-generated text from a large-language model or likely AI-generated text that was likely revised using an AI-paraphrase tool or word spinner.

False positives (incorrectly flagging human-written text as AI-generated) are a possibility in AI models.

AI detection scores under 20%, which we do not surface in new reports, have a higher likelihood of false positives. To reduce the likelihood of misinterpretation, no score or highlights are attributed and are indicated with an asterisk in the report (*%).

The AI writing percentage should not be the sole basis to determine whether misconduct has occurred. The reviewer/instructor should use the percentage as a means to start a formative conversation with their student and/or use it to examine the submitted assignment in accordance with their school's policies.

What does 'qualifying text' mean?

Our model only processes qualifying text in the form of long-form writing. Long-form writing means individual sentences contained in paragraphs that make up a longer piece of written work, such as an essay, a dissertation, or an article, etc. Qualifying text that has been determined to be likely AI-generated will be highlighted in cyan in the submission, and likely AI-generated and then likely AI-paraphrased will be highlighted purple.

Non-qualifying text, such as bullet points, annotated bibliographies, etc., will not be processed and can create disparity between the submission highlights and the percentage shown.